

UNIVERSITY OF BERGEN
DEPARTMENT OF INFORMATICS

Exploring Hardware Agnostic Multiarrays in Magnolia

Author: Marius Kleppe Larnøy

Supervisors: Magne Haveraaen



Bergen
Language
Design
Laboratory

August, 2022

Abstract

We present a specification and implementation of a generic multiarray API based on A Mathematics of Arrays in the general purpose research language Magnolia. We show how we can lift the reasoning on arrays to a more abstract level, and how this enables us to precisely manipulate arrays independent of hardware memory layouts.

Acknowledgements

This thesis has benefited from the Experimental Infrastructure for Exploration of Exascale Computing (eX³), which is financially supported by the Research Council of Norway under contract 270053.

I want to thank my supervisor Magne Haveraaen for his guidance on this thesis, and for facilitating an environment for interesting discussions on generic programming and formal methods through both the Magnolia work group and BLDL. Thank you to my colleagues, and especially Benjamin Chetioui for endless feedback and support on my work. Lastly I want to thank my family for continued support throughout this process.

Marius Kleppe Larnøy
Tuesday 2nd August, 2022

Contents

1	Introduction	2
1.1	Motivations	2
1.2	Contribution	3
1.3	Thesis Outline	3
2	A Mathematics of Arrays	5
2.1	Introduction	5
2.2	The Original MoA Approach	5
2.2.1	Terminology	6
2.2.2	Unary Operations for Array Shapes	6
2.2.3	Indexing	7
2.2.4	Psi indexing	7
2.2.5	Take and Drop	8
2.2.6	Catenate	9
2.2.7	Array Transformations	10
2.2.8	Denotational Normal Form	11
2.2.9	Operational Normal Form	12
2.3	BLDL Approach	14
2.3.1	Canonical rewrite system	14
2.3.2	Padding	14
2.4	Reflection on the approaches	15
3	Magnolia	16
3.1	The Magnolia Language	16
3.2	Related works on Magnolia	19
4	Arrays in other programming languages & PyWake	21
4.1	Arrays in programming languages	21
4.1.1	C	21

4.1.2	Fortran	22
4.1.3	Python	23
4.2	Example domain: wind farm modelling with PyWake	25
5	MoA in Magnolia	28
5.1	Specification	28
5.2	Implementation	30
5.3	Summary	35
5.4	Related works on MoA implementations	36
6	Array Optimizations	37
6.1	P ³ Problem and Magnolia language: Specializing Array Computations for Emerging Architectures	37
6.2	A CUDA implementation	68
7	Future Work & Conclusion	72
7.1	Future Work	72
7.2	Conclusion	73
	Glossary	74
	Bibliography	76

List of Figures

4.1	Birds eye view of the wind farm, AEP of each turbine	27
6.1	Dataflow between CPU(blue) and GPU(green), 1st iteration of the solver.	69
6.2	Improved dataflow between CPU(blue) and GPU(green)	70

List of Tables

6.1	10 steps of the CUDA PDE solver with array dimensions 512^3 , ran on a Nvidia Volta A100/80GB. Timed in bash with <code>time</code>	69
6.2	Snippet of <code>gpumemtimesum</code> result generated from <code>nsys profile <pde.bin></code>	70

Listings

3.1	Concept of a semigroup	17
3.2	Concept of an abelian monoid	18
3.3	Concept of a semiring	18
3.4	External implementation and program in Magnolia	19
3.5	Asserting a claim that a program models a concept	19
3.6	Example output of the Natural Numbers program	19
4.1	C array example	22
4.2	Assembly output of C example	22
4.3	Fortran90 array access example	23
4.4	Assembly output of Fortran example	23
4.5	NumPy basic operations	24
4.6	PyWake example	26
5.1	MoA Signature	29
5.2	MoA Axioms	29
5.3	Signature for mapped operations	30
5.4	Axioms for mapped operations	30
5.5	External While-loop in Magnolia with 1 obs variable and 1 upd variable.	31
5.6	Snippet of array externals in C++	31
5.7	Array externals in Magnolia	32
5.8	Backend definition of a Float64 type with arithmetic operations	32
5.9	Magnolia definition of a Float64 type with arithmetic operations	33
5.10	Implementation of cat in Magnolia	33
5.11	Implementation of take and drop in Magnolia	34
5.12	Implementation of MoA transformations in Magnolia	34
5.13	Array program parameterized with a Float64 element type.	35
6.1	CUDA Annotations	68

Chapter 1

Introduction

Gordon Moore postulated in the 1960s that the number of transistors in a processing unit would double every two years [14]. This postulate largely holds true as we enter the 2020s, with computing power reaching exascale levels (10^{18} FLOPS) in 2018. As hardware continues to evolve we are reliant on software capable of adapting to both current and future architectures, whilst remaining maintainable.

Fields of both research and industry that deal with large volumes of data often utilize HPC – i.e. supercomputers or clusters – to process and perform calculations efficiently. This creates the need for software capable of leveraging distributed architectures, whilst remaining maintainable. MoA [40] is a calculus for working with arrays, generalizing the notion of an array to the concepts of shapes and dimensions. A big motivation behind creating this calculus was how arrays are mapped down to hardware, and how we can rearrange and manipulate the arrays independently of memory layout without losing the ability to target specific architectures.

In this thesis we will explore the MoA calculus, using the generic programming language Magnolia [2, 5] as our vehicle to implement a generic array API based on MoA. We will observe how MoA allows us to manipulate arrays on a hardware independent level without compromising neither performance or supported hardware.

1.1 Motivations

There are numerous domains in science and industry that rely on discretizations of formally defined physics models. These models are reliant on numerical solutions in order

to be applicable to real-world problems, and discretization of PDEs in order to be able to compute finite solutions. A good example is the Navier-Stokes equations, which are used to model the behavior of fluids, e.g. wind and water. In fields such as meteorology and industry sectors such as wind farms, fast and accurate modeling of wind and water is essential. Work by BLDL at the University of Bergen has contributed to the idea that mapping array expressions that can run efficiently on arbitrary hardware is worth exploring. A recurring case study from this research is how to use arrays to compute numerical solutions to PDEs on different hardware, potentially bringing together domain experts in fields such as meteorology who are looking for both faster and more portable ways to simulate data independent of current computing power.

1.2 Contribution

This thesis is comprised out of existing theory in conjunction with the authors own work on applications. Mainly, A Mathematics of Arrays [40] is the work of Lenore Mullin, with subsequent publications on MoA being the work of Mullin and her co-authors. More recent literature on MoA [6, 7, 9] is the work of researchers associated with BLDL in collaboration with Mullin. Additionally, Magnolia is a research language under active development at BLDL. Among the related works are two compiler implementations [2, 5].

What this thesis aims to do is to explore the MoA calculus through the lens of formal specifications. We establish a baseline of understanding of multiarrays and Magnolia, and then we present a specification and implementation for a subset of MoA.

1.3 Thesis Outline

The thesis is structured as follows:

- Chapter 2 introduces the relevant parts of the MoA theory,
- Chapter 3 introduces the Magnolia programming language
- Chapter 4 explores arrays in conventional programming languages, and we motivate the problem at hand by highlighting a relevant domain where efficient array computations are important,

- Chapter 5 describes an implementation of a subset of MoA in the Magnolia programming language,
- Chapter 6 consists of a collaborative article highlighting current work on being carried out at BLDL, and a closer look at an experimental implementation in CUDA,
- Chapter 7 rounds off the thesis with discussion and reflection on both the work that has been done and future work.

Chapter 2

A Mathematics of Arrays

2.1 Introduction

MoA is a theoretical framework for multidimensional arrays, defining them by the notion of their shape. Its inception was the PhD thesis of Lenore Mullin in 1988 [40], and she has been the driving force behind promoting MoA and its applications in parallel computing [16], HPC [15] etc.

The original presentation of MoA draws heavy inspiration from Ken Iverson and APL [25], both in its approach to defining arrays and its notational style. More recent publications depart from the APL roots of the theory [1, 6, 7], focusing instead on its ideas of creating dense array expressions operating on single multiarrays. In this chapter we present an overview of the core theory, first as presented in the original papers and then as given in recent literature. We then draw comparisons between the two approaches.

2.2 The Original MoA Approach

Following from APL where the centerpiece data type is the multidimensional array, MoA revolves around a single array type. Unary operators and infix binary operations are used without any specified operator precedence, and parentheses are used to dictate order of application. Expressions associate implicitly to the right, following from APL.

2.2.1 Terminology

Every array has a *dimensionality*, often denoted by an integer superscript. E.g. ξ^3 is a 3-dimensional array. An array's *shape* denotes the length of each of its dimensions, collected in a 1-dimensional array. E.g. ξ^3 would have a shape of the form $\langle i \ j \ k \rangle$, where $i, j, k \in \mathbb{N}^+$.

Some arrays with a specific dimensionality are more commonly used than others, and as such they are given unique names.

- A *scalar* refers to a 0-dimensional array, i.e. an array with zero dimensions and an empty shape. The literature denotes the empty scalar as σ .
- A *vector* in MoA is a 1-dimensional array of elements with ordered integer indices. In the literature, the empty vector is denoted Θ , and vectors in general are denoted with the typical arrow notation. E.g. $\vec{v} = \langle 1 \ 2 \ 3 \rangle$ is a vector with 3 elements.
- A *matrix* is a 2-dimensional array.

2.2.2 Unary Operations for Array Shapes

- δ (Delta): takes an array and returns its dimensions as a scalar. For scalar arguments $\delta\sigma = 0$.
- ρ (Rho): takes an array and returns its shape as a vector. In particular $\rho\sigma = \Theta$.
- τ (Tau): takes an array returns the total number of objects in the array as a scalar. $\tau\sigma = 1$.

Example: if we have a vector $\vec{v} = \langle 1 \ 2 \ 3 \rangle$, $\tau\vec{v} = 3$.

- ι (iota): takes as argument a scalar $\sigma \in \mathbb{N}$ and generates a vector containing the integer sequence $0 \dots (\sigma - 1)$. For $\sigma = 0$, $\iota 0 \equiv \Theta$.

2.2.3 Indexing

Indexing can be done a few different ways. $\vec{v} = \langle 1 \ 2 \ 3 \rangle$ $\tau \vec{v} = 3$

1. Scalar indexing

$$\vec{v}[0] = 1$$

$$\vec{v}[1] = 2$$

$$\vec{v}[2] = 3$$

2. Vector indexing, given that the components of the index vector all are valid indices for the indexed vector

$$\vec{u} = \langle 2 \ 1 \ 0 \rangle$$

$$\tau \vec{u} = 3$$

$$\vec{v}[\vec{u}] = \langle 3 \ 2 \ 1 \rangle$$

2.2.4 Psi indexing

The psi indexing function ψ takes as a left argument an index vector and right argument an n -dimensional array.

We will begin defining the few special cases for ψ , and then move on to the general form. For the empty scalar σ and the empty vector Θ acts as neutral elements:

$$\Theta \ \psi \ \sigma \equiv \sigma$$

$$\Theta \ \psi \ \vec{x} \equiv \vec{x}$$

$$\Theta \ \psi \ \zeta^n \equiv \zeta^n$$

$$0 \leq i < (\tau \vec{x}) \quad \langle i \rangle \ \psi \ \vec{x} \equiv \vec{x}[i]$$

In general, for an index vector \vec{i} satisfying the bounds¹ $0 \leq^* \vec{i} <^* (\rho \zeta^n)$:

$$\vec{i} \ \psi \ \zeta^n = \zeta^n[\vec{i}[0]; \dots; \vec{i}[n-1]]$$

¹In her dissertation, Mullin introduces the notation $\xi_l \ R^* \ \xi_r$ to talk about constraints on indices. It is used to express the condition that \vec{i} is a valid index vector. E.g. $0 \leq^* \vec{i} <^* (\rho \zeta^n)$.

Partial indexing

Until now we have assumed indexing arguments to ψ to be total, i.e. $(\tau \vec{i}) = \delta \zeta^n = n$. When \vec{i} is a total index, $\vec{i} \psi \zeta^n$ will return a scalar with an empty shape. Now we are going to introduce a partial index, also called a *short* index, i.e. an index vector that does not access precisely to the scalar level, but rather to a subarray level. We impose some restrictions on the partial index vectors, to make it play nicely in bounds of the accessed arrays.

For an index vector \vec{j} and a n -dimensional array ξ^n :

$$\begin{aligned} \text{Given } 0 \leq^* \vec{j} <^* ((\tau \vec{j}) \uparrow (\rho \xi^n)) \\ \text{then } 0 \leq (\tau \vec{j}) \leq \delta(\xi^n) \end{aligned}$$

When \vec{j} is a partial index, the shape of the indexed subarray is given as $\rho(\vec{j} \psi \xi^n) = (\tau \vec{j}) \downarrow (\rho \xi^n)$.

2.2.5 Take and Drop

We can define the \uparrow (*take*) and \downarrow (*drop*) operators as shorthand for indexing on the primary axis². By *take* we are accessing the first n subarrays of the given array, keeping them. Conversely, *drop* will ignore the n first subarrays and keep the rest in an array. When applied to a 1-dimensional array, *take* and *drop* accesses down to the element level.

Example (1 dimensional):

$$\begin{aligned} \vec{x} &= \langle 1 \ 2 \ 3 \ 4 \ 5 \ 6 \rangle \\ \tau \vec{x} &= 6 \\ \vec{x}[\langle 0 \ 1 \ 2 \rangle] &= 3 \uparrow \vec{x} = -3 \downarrow \vec{x} = \langle 1 \ 2 \ 3 \rangle \\ \vec{x}[\langle 4 \ 5 \rangle] &= -2 \uparrow \vec{x} = 4 \downarrow \vec{x} = \langle 5 \ 6 \rangle \end{aligned}$$

²In the literature the theory is usually presented in row-major fashion, and as such when we refer to the primary axis we are talking about the outmost row-axis unless specified otherwise.

Example ($\delta(A) > 1$):

$$A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix}$$

$$0 \uparrow A = \langle 1 \ 2 \ 3 \ 4 \rangle$$

$$0 \downarrow A = \begin{bmatrix} 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix}$$

2.2.6 Catenate

Given two vectors \vec{x} and \vec{y} , their (con)catenation $\vec{x} \# \vec{y}$ yields a vector by indexing them together. The resulting vector $\tau(\vec{x} \# \vec{y}) \equiv (\tau\vec{x}) + (\tau\vec{y})$.

Catenating a vector with a scalar is legal, "promoting" the scalar to a one-element vector.

$$\tau(\vec{x} \# \sigma) \equiv (\tau\vec{x}) + 1$$

Catenation of arrays

Arrays can also be catenated on the primary axis, given that the shapes of the rest of the dimensions match. The shape of two catenated arrays is

$$((1 \uparrow (\rho\xi^n)) + (1 \uparrow (\rho\zeta^n))) \# (1 \downarrow (\rho\zeta^n))$$

Example:

$$A = \begin{bmatrix} 1 & 2 \\ 5 & 6 \\ 9 & 10 \end{bmatrix} \quad B = \begin{bmatrix} 3 & 4 \\ 7 & 8 \\ 11 & 12 \end{bmatrix}$$

$$A \# B = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix}$$

2.2.7 Array Transformations

Here we will introduce three operations for manipulating the indexing of existing arrays: reverse, rotate and transpose.

Reverse

The unary reverse operation ϕ takes an array argument and reverses the order of the elements on the primary axis. It does not change the shape of the array.

$$\rho(\phi\xi^n) \equiv \rho\xi^n$$

For valid indices $0 \leq i < (\rho\xi^n)[0]$:

$$\langle i \rangle \psi (\phi\xi^n) \equiv \langle (\rho\xi^n)[0] - (i + 1) \rangle \psi \xi^n$$

Example:

$$\begin{aligned} \rho A &= \langle 3 \ 4 \rangle & \rho(\phi A) &= \langle 3 \ 4 \rangle = \rho A \\ A &= \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} & \phi A &= \begin{bmatrix} 9 & 10 & 11 & 12 \\ 5 & 6 & 7 & 8 \\ 1 & 2 & 3 & 4 \end{bmatrix} \end{aligned}$$

Rotate

Rotate - \ominus - takes a scalar left argument and an array on the right. $\sigma \ominus \xi^n$ shifts the order of the elements on the primary axis by σ . $\rho(\sigma \ominus \xi^n) \equiv \rho\xi^n$.

$$\sigma \ominus \xi^n \equiv \begin{cases} (\sigma \downarrow \xi^n) \# (\sigma \uparrow \xi^n), & 0 < \sigma \leq (\rho\xi^n)[0] \\ (\sigma \uparrow \xi^n) \# (\sigma \downarrow \xi^n), & -(\rho\xi^n)[0] \leq \sigma < 0 \end{cases}$$

Example:

$$\begin{aligned} \rho A &= \langle 3 \ 4 \rangle & \rho(1 \ominus A) &= \langle 3 \ 4 \rangle = \rho A \\ A &= \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} & 1 \ominus A &= \begin{bmatrix} 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 1 & 2 & 3 & 4 \end{bmatrix} \end{aligned}$$

Transpose

Transpose - \mathbb{Q} - reverses the order of the indices. The resulting shape of the transposed array is the reversal of the original shape.

$$\rho(\mathbb{Q}\xi^n) \equiv \phi(\rho\xi^n)$$

For valid index vectors $0 \leq^* \vec{i} <^* \phi(\rho\xi^n)$,

$$\vec{i} \psi (\mathbb{Q}\xi^n) \equiv (\phi\vec{i}) \psi \xi^n$$

Example:

$$\begin{aligned} \rho A = \langle 3 \ 4 \rangle \quad \rho(\mathbb{Q}A) = \langle 4 \ 3 \rangle = \phi(\rho A) \\ A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} \quad \mathbb{Q}A = \begin{bmatrix} 1 & 5 & 9 \\ 2 & 6 & 10 \\ 3 & 7 & 11 \\ 4 & 8 & 12 \end{bmatrix} \end{aligned}$$

2.2.8 Denotational Normal Form

DNF is a semantics-only, layout-agnostic normal form for array expressions, and represents the most "cost-effective" way to perform operations on any given array. Any MoA expression can be rewritten as a DNF expression. What we want to achieve by reducing an expression to its corresponding DNF is to end up with (ideally) only terms utilizing the ψ operator, which in essence is an array access (cheap). By layout-agnostic we mean that this reduction to normal form is performed *before* any mapping to hardware takes place, meaning optimizations for specific hardware architectures will take place later, and that we don't have to worry about different layouts when performing the conversion to DNF.

ψ -reduction

ψ -reduction is the formal process of transforming an array expression stepwise into its normal form. It is a mechanical process where each operator has its defined rewrite rules. The complete list of rewrite rules were defined and published already in the 1990s, including that ψ -reduction could be performed by a computer [33]. It is the first step in using MoA for efficient array calculations.

In this section we will be discussing the semantics of ψ -reduction. The paper presented in Chapter 6 applies ψ -reduction as part of optimizing array expressions.

Shape Analysis

In order for us to perform a reduction, the shape of both the initial variables of the array expression and the partial results needs to be computed. We also need to define the valid indices (or index vectors) of the result. The valid indices for the expression can be described using:

$$0 \leq i < (\tau E) \text{ for scalar indices or } 0 \leq^* \vec{i} \leq^* (\rho E) \text{ for index vectors}$$

Reduction

Using valid indices as left side argument of ψ , we start out on the form

$$\langle i \rangle \psi E \text{ for scalar indices or } \vec{i} \psi E \text{ for index vectors}$$

The reduction is performed simply by applying what definitions/properties/identities that apply to the given expression. Mullin and Thibault defines a complete list of reduction rules, and shows that the reduction process is deterministic.

2.2.9 Operational Normal Form

Memory layouts in hardware are linear. What separates different systems are how they are accessed, with factors such as the number of processors and processor cores also affecting the final physical layout. Here we will introduce the ONF, a collection of functions to both transform and work with array expressions in a hardware specific context. This part of the MoA algebra is not in the scope of this thesis, but an overview is provided for completeness.

Ravel

Ravel – *rav* – takes an array expression and returns it as a vector, flattening its elements into a one dimensional array. *rav* can be performed both row-major and column-major, depending on the target memory layout.

Reshape

Reshape - \hat{p} . Takes a shape s and an array A and returns A with the same elements but with the shape s . I.e. $\rho(\text{reshape}(s, A)) = s$.

Gamma

To be able to map the DNF efficiently and optimally down to ONF we need to know two things. We need knowledge about the memory layout where our array is being mapped to, and we need an array expression in DNF form. We now introduce a **family** of functions γ which we can use to express the relation between an index and an offset in flat memory.

For vectors, the γ function is layout independent, as vectors are already one-dimensional and contiguous:

$$\gamma(\langle i \rangle, \vec{v}) \equiv \vec{v}[i]$$

The γ functions for n -dimensional arrays depend on the target memory layout. We will give an example on how one can define a γ function for a row-major architecture. Given an n -dimensional array ζ^n with shape $\rho(\zeta^n) = \langle s_0 \dots s_{n-1} \rangle$ and valid index vectors $0 \leq^* \vec{i} <^* \rho\zeta^n$, we can give the following relation:

$$\gamma_{\text{row}}(\vec{i}, \rho(\zeta^n)) = \gamma_{\text{row}}(\langle i_0 \dots i_{n-1} \rangle, \langle s_0 \dots s_{n-1} \rangle) \equiv i_{n-1} + s_{n-1} \times \gamma_{\text{row}}(\langle i_0 \dots i_{n-2} \rangle, \langle s_0 \dots s_{n-2} \rangle)$$

We are now equipped with the operations we need to manipulate arrays to fit specific memory layouts. This concludes our small introduction to the original MoA formalism.

2.3 BLDL Approach

For the last few years, effort has been put into creating a complete pipeline for using MoA in tandem with the Magnolia programming language to generate high-performing array expressions. This work has been carried out by researchers at BLDL in collaboration with Mullin, resulting in a series of papers [6, 7, 9]. A case study on creating an efficient PDE solver serves as the recurring domain of interest. The work includes important additions to the existing theory, as well as required proofs for existing theory.

Here we will briefly introduce this approach to the theory. As the two approaches constitute the same theory, we will focus on the specific contributions of the articles.

2.3.1 Canonical rewrite system

Chetioui et al. approaches MoA with a specific goal in mind, namely to define the minimal array API described in Burrows et al. in terms of the MoA formalism. A subset of MoA is sufficient to investigate this API, and the paper provides rewrite rules for these operations, along with proofs that the rules form a canonical rewrite system. That is, the system is confluent, and any expression can be rewritten to its normal form in a finite number of steps.

2.3.2 Padding

Padding in MoA was introduced by Chetioui et al. as a way to introduce redundancy in arrays. When working with high-performance computers, a limiting factor for computational efficiency is data locality. Large distributed systems might not have sufficient global memory, relying on message passing systems such as MPI to transfer data between components running parallel computations. By prepending or appending existing data to the array, one can limit the need for interaction between different processors. Chetioui et al. demonstrated significant runtime improvements for the PDE solver by padding the arrays.

2.4 Reflection on the approaches

Mullin was heavily inspired by APL when developing the original MoA theory. This is reflected in both in choice of syntax, as well as semantics. Notation used in MoA as presented by Mullin often follows directly from APL. Some semantic rules carry over as well, such as operators being implicitly associated to the right if no parentheses are provided.

The contemporary approach makes an effort to step away from the APL roots. While faithful and equivalent to the original theory, steps are made to create a more seamless transition into a Magnolia flavoured notational style. All operations are now defined consistently throughout in terms of the ψ operator, on the form *index* ψ *op(args)*, leaving the recursive definitions behind. There exist multiple reasons to why this notational style could be preferable.

1. The notational style is well suited for application in parallel computing. By describing the result at each index, the computation can easily be distributed between different processes.
2. Magnolia explicitly disallows recursion, so by relying on recursive definitions there would be a notational gap between implementations and the underlying theory.
3. The recent publications are focused on a limited subset of the theory, keeping notation consistent is more practical.

The approaches are also shaped by their respective goals. In the introduction of her thesis, Mullin argues that MoA enables verification of computer architecture design, building on VLSI design, i.e. design of integrated circuits. The focus of the papers produced at BLDL has been array transformations. Designing a pipeline for creating dense, hardware-independent array expressions in DNF which then can be translated to padded hardware-specific expressions well suited for distributed computing. Combined with Magnolia this creates a platform for theoretically well founded, portable array code not limited to existing hardware.

With these areas of applications in mind, it highlights the versatility of MoA as a theoretical framework for multiarrays.

Chapter 3

Magnolia

Magnolia is a research programming and specification language based on institution theory [13], with the goal of fully capturing Stepanov-style generics [10]. This type of generics is known as *genericity by property* in terms of the Gibbons taxonomy [12], which describes structures and algorithms in terms of syntactic and semantic requirements. To describe the type of genericity provided by Magnolia, Chetioui et al. coined the term *genericity by host language*. Magnolia is dependent on being parameterized by a host language because it provides no base data types or data structures, giving rise to a minor distinction.

Here we will give an introduction to the Magnolia specification and programming language.

3.1 The Magnolia Language

In Magnolia there are four top level module types. A **signature** is a collection of generic type - and function names. We can equip a signature with **axioms**, stating behavior of the defined types and functions. Signatures with axioms together form a **concept**. The **implementation** module expands on the functionality of the signature by allowing us to provide generic implementations for the defined types and functions. This is done either by providing an implementation in a backend language of choice¹, or by providing bodies to the defined functions. Types and structures expecting an external implementation are prefixed with the **require** keyword. In addition to functions, Magnolia also supports

¹At the time of writing this thesis, C++ and Python [8] are supported backend languages

predicates and procedures. Whereas functions are immutable, procedures can modify state of its input variables depending on its mode. Input parameters to procedures must be given explicit mode declarations, which can be either **obs**, **upd** or **out**. An **implementation** fully parameterized by the backend is called a **program** module.

Introductory Example²

Let us look at how to specify and implement natural numbers in Magnolia as an example. We will take advantage of the fact that natural numbers with addition and multiplication form a commutative semiring.

A commutative semiring is defined as a set S with two binary operations *plus* and *mult* such that:

- $(S, plus)$ is a commutative monoid with identity element 0
- $(S, mult)$ is a commutative monoid with identity element 1
- *mult* distributes over *plus*
- *mult* by 0 annihilates S

First, we will define a generic commutative semiring and then we will relate it to a specific natural number implementation. This also serves as an example to show how separating generic structures from specific implementation can decrease code duplication by reusing generic code.

We will begin by specifying a basic algebraic structure, the semigroup. A semigroup is a type together with an associative binary operation. This is straight forward to formulate in Magnolia.

```

1 concept Semigroup = {
2
3     type S;
4
5     function bop(s1: S, s2: S): S;
6
7     axiom associative(s1: S, s2: S, s3: S) {
8         assert bop(bop(s1, s2), s3) == bop(s1, bop(s2, s3));
9     }
10 }

```

Listing 3.1: Concept of a semigroup

²The complete example is available online under `examples/naturalnumbers` [30]

We can then expand on our semigroup concept by adding an identity element, this gives us a monoid. For a semiring to be commutative we require that its underlying monoids are commutative. By adding the commutative property to our monoid concept, we get a commutative (or abelian) monoid. Magnolia’s powerful *renaming* mechanism is also in use here. Renamings allow us to give new names to any declared type or function in a module, providing great flexibility for both specializing generic structures and avoiding unintended name overlaps when importing multiple modules.

```

1 concept AbelianMonoid = {
2
3   // include Semigroup, rename type S to M
4   use Semigroup[S => M];
5
6   function identity(): M;
7
8   axiom idAxiom(m: M) {
9     assert bop(identity(), m) == m;
10    assert bop(m, identity()) == m;
11  }
12  axiom commutative(m1: M, m2: M) {
13    assert bop(m1, m2) == bop(m2, m1);
14  }
15 }

```

Listing 3.2: Concept of an abelian monoid

We can now specify our semiring by bringing in our monoid concept in scope with the appropriate renamings, and by asserting the distribution- and annihilation properties.

```

1 concept Semiring = {
2   // Gives us + and 0
3   use AbelianMonoid[bop => _+_, identity => zero];
4   // Gives us * and 1
5   use AbelianMonoid[bop => _*__, identity => one];
6
7   // Multiplication distributes over addition
8   axiom multDistribution(m1: M, m2: M, m3: M) {
9     assert m1 * (m2 + m3) == (m1 * m2) + (m1 * m3);
10    assert (m1 + m2) * m3 == (m1 * m3) + (m2 * m3);
11  }
12
13   // Annihilation of mult by zero
14   axiom multAnnihilation(m: M) {
15     assert m * zero() == zero();
16     assert zero() * m == zero();
17   }
18 }

```

Listing 3.3: Concept of a semiring

This concludes our specification of a generic commutative semiring, and we proceed to a concrete implementation. We must rely on externally provided types for our implementation because Magnolia does not provide any concrete types. In this example, we’ll use a C++ backend.

```

1 // externally defined types and functions
2 implementation ExternalNat = external C++ base.nat {
3     type Nat;
4
5     function zero(): Nat;
6     function one(): Nat;
7
8     function add(a: Nat, b: Nat): Nat;
9     function mul(a: Nat, b: Nat): Nat;
10 }
11
12 program NaturalNumbers = {
13     use ExternalNat;
14 }

```

Listing 3.4: External implementation and program in Magnolia

Now we want to relate our generic specification to our concrete implementation. The `satisfaction` construct allows us to do precisely this.

```

1 satisfaction NaturalNumbersModelsSemiring = NaturalNumbers
2     models Semiring[M => Nat,
3         zero => zero,
4         one  => one,
5         _+_  => add,
6         _*_  => mul];

```

Listing 3.5: Asserting a claim that a `program` models a `concept`

This satisfaction relation `NaturalNumbersModelsSemiring` expresses that the program `NaturalNumbers` satisfies the axioms of `Semiring` with the provided type `Nat` and functions `zero`, `one`, `add` and `mul`.

If we provide a minimal backend in C++, together with a main function with some test calls, we can compile with `magnoliac` and check that our small specification and implementation in fact yields executable code.

```

1 $ ./natnum.bin
2 zero(): 0
3 add(one(), one()) = 2

```

Listing 3.6: Example output of the Natural Numbers program

3.2 Related works on Magnolia

There has been a wide range of work done in the Magnolia ecosystem since its inception.

- Bagge lays the groundwork for the concepts explored in Magnolia, as well as providing the first compiler implementation.

- Haugsbakk uses Magnolia as a vehicle to explore program transformation techniques.
- Abusdal explores the MoA calculus in the context of Magnolia similarly to this thesis, but through the lens of array transformations.
- Chetioui et al. highlights a redesigned Magnolia compiler [5], extending it to support a Python backend, and implementing a subset of the Boost Graph Library [42] in Magnolia to demonstrate how this allows for performant code in both transpiled C++ and Python from the same Magnolia source.
- Hamre provides insights on the viability and use of third-party verification software such as SMT solvers to prove or disprove satisfaction claims in Magnolia code.

Chapter 4

Arrays in other programming languages & PyWake

4.1 Arrays in programming languages

Arrays in programming are used to conveniently allocate equal parts of contiguous memory without having to refer to individual variables for each allocation. In this section we will take a short detour to look at multiarray support in some conventional programming languages.

4.1.1 C

C does not have support for multiarrays, only allowing integer indexing. While one can both create and index C arrays with multiple integer indices – e.g. a "2-dimensional" array `int array[4][4]` – this is just syntactic sugar for making 1-dimensional array manipulation easier. Array creation in C is in itself in fact syntactic sugar for creating a pointer to a location in memory. C arrays are in reality pointers to the first block of memory, and accessing elements after that is just providing an offset to the initial pointer position. This is reflected in C11 standard [24], where it describes array memory layout as contiguous. Two arrays of different dimensionality with the same elements in identical order will be represented equally in memory. C also allows for nesting of arrays, i.e. arrays of arrays.

Abusdal showcased how efficient the C compiler can be, in an example similar to this.

```
1 const int one_d[4] = {1,2,3,4};
2 const int two_d[2][3] = {{1,2},{3,4}};
3
4 int c_indexing() {
5     if(one_d[0] == two_d[0][0] &&
6         one_d[1] == two_d[0][1] &&
7         one_d[2] == two_d[1][0] &&
8         one_d[3] == two_d[1][1])
9         {return 5;}
10    else
11        {return 10;}
12 }
```

Listing 4.1: C array example

Compiling using GCC with all optimizations on, we can see that the whole comparison in the `c_indexing` function has been reduced to a single `mov` instruction.

```
1 <c_indexing>:
2     mov $0x5, %eax
3     ret
```

Listing 4.2: Assembly output of C example

4.1.2 Fortran

Fortran arrays offer much more fine-grained manipulations out of the box. It supports arithmetic operations on arrays, reducing the need to iterate through arrays as one are used to in the C school of languages. Interestingly, unlike C, Fortran does not allow for nested arrays. It is explicitly described in the Fortran 2003 standard [23] that *a **scalar** is a datum that is not an array (..) an **array** is a set of scalar data, all of the same type (..)*. Fortran also limits the number of dimensions an array can have to seven, but poses no restrictions on the number of elements each dimension can have.

Drawing once again from Abusdal, the GCC Fortran compiler can – as the C compiler – optimize out all array accesses given the right conditions.

```

1 function fortran_indexing() result(r)
2
3
4 integer :: r
5 integer, dimension(2,2,2) :: array1
6 integer, dimension(8) :: array2
7 array1 = reshape([1,2,3,4,5,6,7,8], [2,2,2])
8 array2 = reshape([1,2,3,4,5,6,7,8], [8])
9
10 if(array2(1) == array1(1,1,1) .and. &
11    array2(2) == array1(2,1,1) .and. &
12    array2(3) == array1(1,2,1) .and. &
13    array2(4) == array1(2,2,1) .and. &
14    array2(5) == array1(1,1,2) .and. &
15    array2(6) == array1(2,1,2) .and. &
16    array2(7) == array1(1,2,2) .and. &
17    array2(8) == array1(2,2,2)) then
18     r = 5
19 else
20     r = 10
21 end if
22 end function access

```

Listing 4.3: Fortran90 array access example

As with 4.1.1, if we compile using GCC with all optimizations enabled and inspect the assembly output we can see that all array accesses has been optimized out, and we are left with a simple mov instruction setting the result to 5.

```

1 <fortran_indexing_>:
2 mov $0x5, %eax
3 ret

```

Listing 4.4: Assembly output of Fortran example

4.1.3 Python

Python is not known for its arrays, but rather for its lists. While arrays typically require a type to be provided, the flexibility of lists are more suited for the dynamically typed paradigm Python follows. Python lists are dynamically sized and allows for mixing of types in a single container. As with other languages, Python's built-in lists has – while useful – largely been replaced with library provided alternatives for performance heavy computations. While the Python standard library includes an array package [11], it is scheduled for deprecation in Python 4, most likely due to third-party libraries already meeting the demand for arrays. Let us take a look at the most prolific one: NumPy.

NumPy [18] is a Python library for array and numerical computations. It has become synonymous with array- and numerical computing in the realm of Python, and is part of the foundation of libraries such as SciPy [46] and pandas [48]. It has even proved itself worthy of application in fields with demanding computational requirements, e.g. astrophysics [35].

NumPy offers a powerful API for array manipulation, and manages to achieve higher performance than usually associated with Python by leveraging optimized C code for much of its core implementation, and Fortran libraries such as OpenBLAS. Additionally, there exists extensions to NumPy designed to leverage high performance architectures such as GPUs [34], backing up the claim that there is a demand for tooling to adapt to increasing computing power. Many of the operations provided by NumPy corresponds with operations from MoA and Fortran, both in naming and behavior. This gives an indication that many array libraries – although different design choices – to a varying degree inherit some core traits from early array languages such as APL and Fortran.

Here we give a few examples of operations that show up in NumPy as well as MoA or Fortran.

```
1 import numpy as np
2 # creating a 1-dimensional array
3 data = np.arange(12)
4 > array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
5
6 # the shape
7 data.shape
8 > (12,)
9
10 # reshape it to an array with shape (3,4)
11 data = data.reshape((3, 4))
12
13 > array([[ 0,  1,  2,  3],
14         [ 4,  5,  6,  7],
15         [ 8,  9, 10, 11]])
16 data.shape
17 > (3,4)
18 # indexing
19 data[1,2] # total
20 > 6
21 data[2] # partial
22 > array([ 8,  9, 10, 11])
23 data[-1] # negative (in bounds)
24 > array([ 8,  9, 10, 11])
25
26 # map
27 data + 2
28 > array([[ 2,  3,  4,  5],
29         [ 6,  7,  8,  9],
30         [10, 11, 12, 13]])
```

Listing 4.5: NumPy basic operations

4.2 Example domain: wind farm modelling with PyWake

Simulating wind in wind farms is an application area well suited for high performance array computations. Here we take a quick look on how the PyWake library combines the computational strength of NumPy with Python's ease of use.

Calculating wind flow is far from a new field of research. The Navier-Stokes equations on the motion of viscous fluids were formulated in the first half of the 19th century. Proofs of general solutions and their uniqueness are famously one of the open questions presented as part of the Millennium Prize Problems [4]. Leveraging the power of computers to calculate wind flow proved useful, Veers exemplifies earlier application areas of this, using simulation data to analyse the aerodynamics of wind turbines.

Arrays were an obvious representation of data to explore when it came to fluid simulations. The three spatial axes of the real world can be represented by a three-dimensional array, and the problem could then be reduced to efficiently propagating updated values between a finite collection of grid points. One also saw the potential for running much of the computation in parallel, especially since the rise of GPGPU in the mid- 2000s [32, 26] with platforms such as CUDA [36] and OpenCL [43].

Large-scale wind farm design and maintenance necessitate extensive modeling of wind and ocean conditions. The wind conditions of the locations suitable for power generation are self-explanatory: you would want to build your wind farm in a location with average wind speeds high enough to generate a sufficient amount of electricity. The addition of wind turbines complicates matters considerably. The rows of turbines disrupt the wind flow, so positioning the turbines to maximize output in the face of disrupted air flow is critical.

PyWake [38] is a Python library developed and maintained by the Technical University of Denmark, used for calculating wind fields and energy production of wind farms. As an academic effort [27, 47, 44, 41], PyWake is a powerful tool capable of simulating wind conditions in wind farms on both sea (accounting for waves) and land (accounting for terrain). PyWake uses XArray [21] for its data representation, which are labeled multiarrays. XArray is in itself built on top of NumPy arrays. With NumPy arrays

serving as the underlying data structures, calculating how wind propagates through the wind farm can be reduced to numerical maps on arrays.

While highly customizable to meet the requirements of domain experts, PyWake comes with a library of pre-defined models to work with. This includes real-world wind turbines that are currently in use, as well as a collection of existing wind farms with which one can experiment. PyWake is also tightly integrated with matplotlib [22], allowing for easy visualization of aspects such as AEP, wind speeds across a wind farm, and wind speeds around single turbines.

In this example we will be using the pre-defined site Horns Rev 1, which is an offshore wind farm in Denmark. We will create a `SimulationResult`, which is a labeled `XArray` containing arrays of information such as wind speeds, wind direction and power production across the site. By invoking methods on the `SimulationResult`, we can easily extract and visualize information, e.g. AEP.

```
1 import numpy as np
2 import py_wake
3
4 from py_wake.examples.data.hornsrev1 import Hornsrev1Site, V80, wt_x,
   ↪ wt_y, wt16_x, wt16_y
5 from py_wake import NOJ
6
7 # selecting type of turbine
8 windTurbines = V80()
9 # selecting site
10 site = Hornsrev1Site()
11 # NOJ is a wake model, which combines a site with a set of turbines
12 noj = NOJ(site, windTurbines)
13
14 # creating a simulation result with 16 turbines
15 sim_res = noj(wt16_x, wt16_y)
```

Listing 4.6: PyWake example

Here we give an example plot to showcase how one can easily present information of about the wind farm. Figure 4.2 gives clear information about how the power production of the inner turbines are being affected by the surrounding ones.




Figure 4.1: Birds eye view of the wind farm, AEP of each turbine

Chapter 5

MoA in Magnolia

In this chapter, we present an implementation of a subset of MoA in the Magnolia programming language, utilizing the `magnoliac` [5] compiler currently under active development. We leverage a C++ backend to provide us with the basic types and structures we need. The complete code base for the implementation presented in this chapter is publically available online [31].

Remark: This implementation is a result of work done in preparation for the paper presented Chapter 6, and reflects its intended use as a platform to explore the API presented in Burrows et al. and Chetioui et al.. As such, it is not a complete implementation of the ψ -calculus, but rather a subset.

Remark 2: Following and release of the previous Magnolia compiler [2], a large standard library was developed. `magnoliac` is at the time of writing this thesis incompatible with the standard library, and as such all modules used in this project have been developed independently of previous work.

5.1 Specification

Everything in the ψ -calculus revolves around a single type: the array. Vectors are 1-dimensional arrays, as are shapes and index vectors. Separating the concepts of array, shape, and index will be useful for our purposes. This is primarily due to the fact that operations (both unary and binary) are defined on legal ranges of shapes, and using the type system to constrain the functions we define will make the specification more readable and less error-prone.

```

1 concept ArrayBaseOps {
2   // Array type
3   type Array;
4   // Index type
5   type Index;
6   // Shape type
7   type Shape;
8   // Element type
9   require type Element;
10
11  /*
12  We separate the integer types based on
13  intended use to avoid type mixups
14  */
15  type Axis;
16  type Dim;
17  type Offset;
18  type Size;
19
20  function dim(a: Array): Dim;
21  function shape(a: Array): Shape;
22  function total(a: Array): Size;
23
24  // Predicates assuring that index-parameters are of the correct size
25  predicate isPartialIndex(i: Index, a: Array);
26  predicate isTotalIndex(i: Index, a: Array);
27
28  function psi(i: Index, a: Array): Array guard isPartialIndex(i, a);
29  function psi(i: Index, a: Array): Element guard isTotalIndex(i, a);
30
31  function cat(a1: Array, a2: Array): Array
32    guard drop(0, shape(a1)) == drop(0, shape(a2));
33
34  function take(o: Offset, a: Array): Array;
35  function drop(o: Offset, a: Array): Array;
36
37  // transformations
38  procedure rotate(obs ax: Axis, obs j: Offset, upd a: Array)
39    guard ax < dim(a);
40  procedure reverse(upd a: Array);
41  procedure transpose(upd a: Array);
42 }

```

Listing 5.1: MoA Signature

Now that the core signature is defined, we can equip it with axioms to assert behavior:

```

1 // rotate does not change the shape of the array
2 axiom rotateShapeAxiom(ax: Axis, j: Offset, a: Array) {
3   var pre_shape = shape(a);
4   call rotate(ax, j, a);
5   assert shape(a) == pre_shape;
6 }
7
8 // transposing an array reverses its shape
9 axiom transposeShapeAxiom(a: Array) {
10  var pre_shape = shape(a);
11  call transpose(a);
12  assert shape(a) == reverse(pre_shape);
13 }

```

Listing 5.2: MoA Axioms

The API described in Burrows et al. states that our implementation is going to need mapped operations on arrays. Let us express this in our specification:

```

1 concept MappedOps = {
2
3   use ArrayBaseOps;
4
5   // Requiring functions that will be provided by the backend
6   require function _+(a: Element, b: Element): Element;
7   require function _-(a: Element, b: Element): Element;
8   require function _*(a: Element, b: Element): Element;
9   require function _/(a: Element, b: Element): Element;
10  require function _-(a: Element): Element;
11
12  require predicate _<_(a: Element, b: Element);
13  require predicate _==_(a: Element, b: Element);
14
15  // Array-Array operations
16  function _+(a: Array, b: Array): Array;
17  function _-(a: Array, b: Array): Array;
18  function _*(a: Array, b: Array): Array;
19  function _/(a: Array, b: Array): Array;
20  function _-(a: Array): Array;
21
22  predicate _==_(a: Array, b: Array);
23
24  // Scalar-Array operations
25  function _+(a: Element, b: Array): Array;
26  function _-(a: Element, b: Array): Array;
27  function _*(a: Element, b: Array): Array;
28  function _/(a: Element, b: Array): Array;
29 }

```

Listing 5.3: Signature for mapped operations

With the MappedOps signature defined, we can add axioms to express the intended semantics and relate the mapped operations to their underlying element-wise operations.

```

1 axiom binaryMap(a: Array, b: Array, ix: Index)
2   guard isTotalIndex(ix, a) && isTotalIndex(ix, b) {
3
4     assert psi(a+b, ix) == psi(a, ix) + psi(b, ix);
5     assert psi(a-b, ix) == psi(a, ix) - psi(b, ix);
6     assert psi(a*b, ix) == psi(a, ix) * psi(b, ix);
7     assert psi(a/b, ix) == psi(a, ix) / psi(b, ix);
8   }
9 axiom scalarLeftMap(e: Element, a: Array, ix: Index)
10   guard isTotalIndex(ix, a) {
11     assert psi(e+a, ix) == e + psi(a, ix);
12     assert psi(e-a, ix) == e - psi(a, ix);
13     assert psi(e*a, ix) == e * psi(a, ix);
14     assert psi(e/a, ix) == e / psi(a, ix);
15   }
16 axiom unaryMap(a: Array, ix: Index) {
17   assert psi(-a, ix) == -psi(a, ix);
18 }

```

Listing 5.4: Axioms for mapped operations

5.2 Implementation

With our specification completed, we can provide an implementation. This is accomplished by combining externally defined functions with our API and providing our own

function and procedure bodies. What we rely on from the backend are:

- A looping mechanism
- An array structure with getters/setters
- Base types

Magnolia does not have built-in support for control structures such as loops, and by design does not allow recursion. While cumbersome, we can implement our own looping structures by leveraging loops present in the backend language at hand. Listing 5.5 is an example of a while-loop with one updatable *state*, and an observable *context*. The `repeat` procedure calls the `body` procedure as long as the `cond` predicate holds true given the context and current state. A simple use case would be printing the elements of a list, providing the list as context, an integer type as state, and `cond` as a upper bound predicate. A drawback to this approach is that Magnolia lacks support for variadics, which forces us to provide different implementations based on the number of contexts and states one would want present in the loop.

```

1 implementation WhileLoop1_1 =
2   external C++ while_loop1_1 {
3     require type Context1;
4     require type State1;
5
6     require predicate cond(context1: Context1, state1: State1);
7     require procedure body(obs context1: Context1,
8                           upd state1: State1);
9     procedure repeat(obs context1: Context1,
10                    upd state1: State1);
11 };

```

Listing 5.5: External While-loop in Magnolia with 1 **obs** variable and 1 **upd** variable.

We define our base types and array data structure in C++, contained in structs.

```

1 template <typename _Element>
2 struct moa {
3   // defining types
4   typedef _Element Element;
5   typedef int Int32;
6   typedef float Float32;
7   typedef std::vector<Int> Index;
8   typedef std::vector<Index> IndexSpace;
9   typedef std::vector<Int> Shape;
10
11   // defining the Array
12   struct Array {
13
14     Element * _content;
15     Shape _shape;
16
17     // total index, returns element
18     inline Element psi(const Index i);
19     // partial index, returns subarray
20     inline Array psi(const Index i);
21 };

```

```

22 // indexing guards
23 inline bool isTotalIndex(const Index i, const Array a)
24     return size(i) == size(a);
25 inline bool isPartialIndex(const Index i, const Array a)
26     return size(i) < size(a);
27 };

```

Listing 5.6: Snippet of array externals in C++

These can then be put together with their corresponding Magnolia-side definitions using the `external` keyword.

```

1 implementation ExternalMoaOps = external C++ moa {
2     require type Element;
3     type Array;
4     type Index;
5     type IndexSpace;
6     type Shape;
7     type Int32;
8     type Float32;
9 }

```

Listing 5.7: Array externals in Magnolia

It is worth noting that the required type `Element` is not defined in our `moa` struct, but is instead passed as a generic template argument. This allows us to pass different `Element` types to our arrays, giving us more flexibility. We provide backend definitions for an `Int64` and a `Float64` element type, as well as operations on the types, in this implementation.

```

1 struct float64_utils
2 {
3     typedef double Float64;
4
5     inline Float64 zero() { return 0.0; }
6     inline Float64 one() { return 1.0; }
7
8     inline Float64 binary_add(const Float64 a, const Float64 b)
9         return a + b;
10    inline Float64 binary_sub(const Float64 a, const Float64 b)
11        return a - b;
12    inline Float64 mul(const Float64 a, const Float64 b)
13        return a * b;
14    inline Float64 div(const Float64 a, const Float64 b)
15        return a / b;
16    inline Float64 unary_sub(const Float64 a)
17        return -a;
18    inline Float64 abs(const Float64 a)
19        return std::abs(a);
20 };

```

Listing 5.8: Backend definition of a `Float64` type with arithmetic operations

We also specify a generic `NumberType`-concept in Magnolia, which in combination with our backend-definition yields our candidates for types we can rename `Element` to:

```

1 concept NumberOps = {
2     type NumberType;
3
4     function zero(): NumberType;
5     function one(): NumberType;
6
7     function binary_add(a: NumberType, b: NumberType): NumberType;
8     function binary_sub(a: NumberType, b: NumberType): NumberType;
9     function mul(a: NumberType, b: NumberType): NumberType;
10    function div(a: NumberType, b: NumberType): NumberType;
11    function unary_sub(a: NumberType): NumberType;
12    function abs(a: NumberType): NumberType;
13 }
14 implementation Float64Utils = external C++ float64_utils
15     NumberOps[NumberType => Float64];

```

Listing 5.9: Magnolia definition of a Float64 type with arithmetic operations

Now we have all the building blocks we need to provide an implementation for our `ArrayBaseOps` concept. We make an effort to reflect the MoA notation introduced in Chetoui et al. For the complete list of definitions used to implement these functions we refer the reader to Chetoui et al., but as an informative example we will compare the definition of catenation to our implementation. The definition given in Chetoui et al. reads:

Given an arrays A and B with $\rho(A) = \langle s_0^A, s_1, \dots, s_{n-1} \rangle$ and $\rho(B) = \langle s_0^B, s_1, \dots, s_{n-1} \rangle$ – i.e. two arrays with identical shape except for the first shape element – we can describe the result at index $\langle i \rangle$ as

$$\langle i \rangle \psi \text{ cat}(A, B) = \begin{cases} \langle i \rangle \psi A & \text{if } i < s_0 \\ \langle i - s_0 \rangle \psi B & \text{otherwise} \end{cases}$$

Notice how $\langle i \rangle$ is a index vector of length 1. For any array C with $\delta(C) > 1$ this is a partial index, and as such we are updating subarrays rather than individual elements, i.e. a map.

We provide a body to a loop with 3 observable contexts and 2 updatable states, and instantiate the loop with a valid index space, a result array with correct dimensions, and an counter variable `c`.

For each valid index, the procedure `cat_ix` is executed once.


```

1 procedure cat_ix(obs a: Array,
2                 obs b: Array,
3                 obs ix: Index,
4                 upd res: Array) {
5
6     var s0 = get(shape(a), zero());
7     var i0 = get(ix, zero());
8
9     if i0 < s0 then {
10         call set(res, ix, psi(ix, a));
11     } else {
12         var new_ix = create_1d_index(i0 - s0);
13         call set(res, ix, psi(new_ix, b));
14     };
15 }

```

Listing 5.10: Implementation of `cat` in Magnolia

All the operations follow this pattern of utilizing an external loop to describe the result at element level or subarray level, depending on whether we have a total or partial index. We continue by providing bodies to our `take` and `drop` prototypes, again calling each procedure once for every valid index.

```

1 procedure take_ix(obs a: Array,
2                 obs ix: Index,
3                 obs t: Offset,
4                 upd res: Array) {
5     if zero() <= t then {
6         call set(res, ix, psi(a, ix));
7     } else {
8         var s0 = get(shape(a), zero());
9         var i0 = get(ix, zero());
10        var new_ix = create_1d_index(s0 - abs(t) + i0);
11        call set(res, ix, psi(new_ix, a));
12    };
13 }
14 procedure drop_ix(obs a: Array,
15                  obs ix: Index,
16                  obs t: Int,
17                  upd res: Array) {
18     if zero() <= t then {
19         var i0 = get(ix, zero());
20         var new_ix = create_1d_index(i0 + t);
21         call set(res, ix, psi(a, new_ix));
22     } else {
23         call set(res, ix, psi(ix, a));
24     };
25 }

```

Listing 5.11: Implementation of `take` and `drop` in Magnolia

With `cat`, `take` and `drop` defined we can implement our transformations, which are defined in terms of these operations [7].

```

1 // reverse
2 procedure reverse_ix(obs a: Array,
3                    obs ix: Index,
4                    upd res: Array) {
5     var sh_0 = get(a, zero());
6     var ix_0 = get(ix, zero());
7

```

```

8 |   var new_ix_0 = sh_0 - (ix_0 + one());
9 |   var new_ix = cat_index(create_1d_index(new_ix_0),
10 |                        drop_index_elem(ix, zero()));
11 |
12 |   call set(res, new_ix, psi(ix, a));
13 | }
14 | // rotate
15 | procedure rotate_ix(obs a: Array,
16 |                   obs ix: Index,
17 |                   obs sigma: Offset,
18 |                   upd res: Array) {
19 |   if zero() <= sigma then {
20 |     var e1 = take(-sigma, psi(ix, a));
21 |     var e2 = drop(-sigma, psi(ix, a));
22 |     call set(res, ix, cat(e1,e2));
23 |   } else {
24 |     var e1 = drop(sigma, psi(ix, a));
25 |     var e2 = take(sigma, psi(ix, a));
26 |     call set(res, ix, cat(e1,e2));
27 |   };
28 | }
29 | // transpose
30 | procedure transpose_ix(obs a: Array,
31 |                      obs ix: Index,
32 |                      upd res: Array) {
33 |   var e = psi(reverse(ix), a);
34 |   call set(res, ix, e);
35 | }

```

Listing 5.12: Implementation of MoA transformations in Magnolia

With transformations complete, we have successfully implemented the API specified by Burrows et al. and Chetoui et al. In order to get executable code, all that remains is to combine our external element type with our MoA implementation in a **program** module.

```

1 | program Float64Arrays = {
2 |   use Float64Utils;
3 |   use MoA[Element => Float64];
4 | }

```

Listing 5.13: Array program parameterized with a Float64 element type.

5.3 Summary

In this chapter we have provided a specification and implementation of a subset of MoA in Magnolia. We have made a case for how separating the generic API from its implementations can provide more flexibility. With the core API defined, the programmer is free to provide any number of different implementations based on domain specific needs, with guarantees of type safety by the compiler.

Utilizing an external loop to serve as the core of our implementation made it possible to describe the result of a computation at each index, closely following the notational style of recent BLDL efforts [6, 7].

While operational, the implementation presented in this chapter has not lived up to its full potential. Crucially, in order to serve as the backbone for the approach presented in section 6.1, a working implementation of circular padding [7] was needed. While effort was put into development, time constraints prevented it from reaching a satisfactory level suitable for use in the computational experiments planned for the publication. This raises an important point. While the Magnolia compiler can provide some type safety guarantees, externally provided code is no less prone to programming errors. The instability in the case of circular padding could be traced back to the C++ backend, and would require a non-trivial rewrite of the indexing code. As such, a separate, more stable implementation ended up being utilized in the article.

With this experience in mind, this brings us back to a remark from the beginning of the chapter. The implementation presented in this chapter was made without the support of the standard library due to incompatibility issues. Modularity in programming allows for greater flexibility, and it stands to reason that a trusted library of modules provides greater assurance that the types and data structures being imported function as intended.

5.4 Related works on MoA implementations

Previous partial or full implementations of MoA exist.

- In 1994, Mullin and Thibault implemented the Psi Compiler [33], demonstrating a working C implementation of ψ -reduction.
- Python MoA [37] is a proof of concept of a Python implementation of MoA funded by Quansight Labs.
- There are currently efforts to implement a MoA library for LFortran [39], headed by Mullin and a group of people connected to the Fortran community.

Chapter 6

Array Optimizations

6.1 P^3 Problem and Magnolia language: Specializing Array Computations for Emerging Architectures

This section consists of a unpublished paper showcasing array transformations using Magnolia axioms to perform rewrites.

The article is at the time of submitting this thesis undergoing peer review. A key aspect we expect to get constructive feedback on is the number of different architectures explored. The CUDA implementation discussed in section 6.2 is a work in progress to address this.

The author of this thesis contributed to the underlying Magnolia implementation of MoA, a CUDA implementation discussed in section 6.2, as well as contributing to parts of the article text.

P³ Problem and Magnolia Language: Specializing Array Computations for Emerging Architectures

Benjamin Chetioui^{1,*}, Marius Kleppe Larnøy¹, Jaakko Järvi², Magne Haveraaen¹, and Lenore Mullin³

¹*Department of Informatics, University of Bergen, Bergen, Norway*

²*Department of Computing, University of Turku, Turku, Finland*

³*College of Engineering and Applied Sciences, University at Albany, SUNY, Albany, USA*

Correspondence*:

Benjamin Chetioui

Institutt for informatikk, Thormøhlens Gate 55, 5008 Bergen, Norway

benjamin.chetioui@uib.no

2 ABSTRACT

3 The problem of producing portable high-performance computing (HPC) software that is cheap
4 to develop and maintain is called the P³ (performance, portability, productivity) problem. Good
5 solutions to the P³ problem have been achieved when the performance profiles of the target
6 machines have been similar. The variety of HPC architectures is, however, large and can be
7 expected to grow larger. Software for HPC therefore needs to be highly adaptable, and there is a
8 pressing need to provide developers with tools to produce software that can target machines with
9 vastly different profiles.

10 Multi-dimensional array manipulation constitutes a core component of numerous numerical
11 methods, such as finite difference solvers of Partial Differential Equations (PDEs). The efficiency
12 of these computations is tightly connected to traversing and distributing array data in a hardware-
13 friendly way. The Mathematics of Arrays (MoA) allows for formally reasoning about array
14 computations and enables systematic transformations of array-based programs, e.g. to use
15 data layouts that fit to a specific architecture.

16 This paper shows a general methodology for solving the P³ problem in a well-specified
17 domain using Magnolia, a language designed to embody generic programming. The Magnolia
18 programmer can restrict the semantic properties of abstract generic types and operations by
19 defining so-called axioms. Axioms can be used to produce tests for concrete implementations of
20 specifications, for formal verification, or to perform semantics-preserving program transformations.

21 We leverage Magnolia's semantic specification facilities to extend the Magnolia compiler with a
22 term rewriting system. We implement MoA's transformation rules in Magnolia, and demonstrate
23 through a case study on a finite difference solver of PDEs how our rewriting system allows
24 exploring the space of possible optimizations.

25 **Keywords:** Partial Differential Equations, Generic Programming, Magnolia Language, Mathematics of Arrays, Term Rewriting, High-
26 Performance Computing

1 INTRODUCTION

The quest for higher performance fuels innovation on hardware architectures; we have seen a wide variety of high-performance computing (HPC) architectures in the past and can expect new ones to keep appearing. Long-lived and successful HPC software must thus be highly adaptable, adjustable to different memory hierarchies and changing intra- and interprocess communication hardware.

The problem of producing portable HPC software that is easy, or at least not unreasonably difficult, to develop and maintain is called the P³ (performance, portability, productivity) problem. Good solutions to the P³ problem have been achieved when the performance profiles of the target machines have been similar (Wolfe, 2021). As more new hardware architectures emerge, there is a pressing need to provide developers with tools to produce such software for targets with vastly different profiles. This includes architectures within Wolfe’s P³ machine performance model (CPUs, GPUs, or other accelerators, possibly distributed) (Wolfe, 2021) but also those that do not (e.g., Groq’s Tensor Streaming Processor (Abts et al., 2020)).

Multidimensional array manipulation is at the core of numerous numerical methods. The topic of optimizing the performance of array computations is therefore extremely relevant to the P³ problem. We have previously explored the Mathematics of Arrays (MoA) formalism (Mullin, 1988) as a tool to optimize array computations for different hardware architectures (Chetioui et al., 2019, 2021). A thorough mathematical understanding of a given domain is key to enabling domain-specific semantic-preserving rewrites — and therefore optimizations.

The portability and productivity pillars of P³ are both strongly related to the notion of code reuse. Portability as meant here is the ability to run the same code with high performance on different machines. Productivity means that applications can be developed and maintained with a reasonable and predictable effort. Research unequivocally shows that productivity increases through reuse (Nazareth and Rothenberger, 2004; Basili et al., 1996; Frakes and Succi, 2001). Generic programming has proven to be an effective method of constructing libraries of reusable software components. The Magnolia programming language (Bagge, 2009) is designed as an embodiment of generic programming (Chetioui et al., 2022). It allows the flexible intermixing of specifications and implementations. Specifications can additionally be restricted by semantic requirements (called *axioms*) in the form of assertions. These axioms can be used for testing (Bagge et al., 2011), but also for optimization when used as directed rewrite rules, in the case of equational or conditional equational axioms (Bagge and Haverlaen, 2009).

1.1 Schedules as Hardware Abstractions

In their 2012 paper on Halide, Ragan-Kelley et al. introduce the term *schedule* to refer to decisions about storage and about the order of computations in a program (Ragan-Kelley et al., 2012). The insight is that the essence of an algorithm is distinct from its schedule — allowing the advent of a programming model where both kinds of computations are not anymore intertwined but instead expressed independently from each other.

Stepanov-style generic programming abstracts algorithms and data structures by specifying minimum syntactic and semantic requirements on instantiations. Said differently, the types and operations underlying a generic implementation are only characterized by the part of their observable behavior that is relevant to the generic algorithm.

When observed through the lens of generic programming, a schedule is an abstraction for the kind of hardware architecture underlying the computations. We consider only the information about the hardware

68 that is relevant for executing our algorithm efficiently: how computations should be ordered, and how data
69 should be stored. Similar hardware architectures are then valid instantiations for the same schedule.

70 Scheduling, in the case of array computations, relates particularly to the access patterns of the arrays.
71 As a motivating example, consider an array program running on a single CPU with memory, the classical
72 model of a computer. We may have three standard traversal patterns for computations over our arrays:

- 73 1. a row-major traversal;
- 74 2. a column-major traversal;
- 75 3. a tiled traversal.

76 While the original algorithm can be expressed without making any assumption about the underlying
77 hardware, the choice of a particular hardware will dictate which traversal pattern is most efficient. In
78 other cases, the choice of a particular schedule may be desirable. E.g., on hardware consisting of several
79 distributed CPUs connected through some communication network, we may want the schedule to handle
80 inter-CPU communication using MPI. If each one of these CPUs is connected to several GPUs, we may
81 also want the schedule to load data on and off the GPUs as needed. Such choices will affect the desired
82 data layout, and consequently the data access patterns so as to match the distribution of the data. These
83 changes will have to be reflected in the presentation of the algorithm.

84 The execution time for an algorithm adapted to its schedule may be dramatically shorter than for an
85 algorithm exhibiting inadapted data access patterns. While an algorithm and its schedule can be expressed
86 independently, choices in the latter may affect what is an appropriate expression of the former, and vice
87 versa. Our approach uses rewriting technology to adapt a unique algorithm to adequately exploit the data
88 traversal pattern of a schedule, and underlying hardware characteristics.

89 Throughout the rest of the paper, we view schedules as hardware abstractions. This view is fully
90 compatible with Ragan-Kelley et al.'s definition of schedules, but conveys our intent more accurately.

91 1.2 Contribution and Structure of the Paper

92 The contribution of this paper is a general methodology for solving the P³ problem in a well-specified
93 domain, that keeps the essence of the algorithm separate from its schedule. We perform a case study on
94 a Partial Differential Equation (PDE) solver based on Finite Difference Methods (FDM). We extend the
95 Magnolia compiler with code generation and term rewriting facilities based on axioms. We implement
96 our solver in Magnolia, using MoA as an underlying basis for the code, giving us both generic and
97 hardware-specific formally verified optimization rules — also directly implemented in Magnolia.

98 The paper is structured as follows. Section 2 provides necessary background on MoA and Magnolia.
99 Section 3 describes our methodology in detail, and illustrates it with a PDE solver based on FDM. Section 4
100 reflects on our work and ties it together with relevant related work.

2 BACKGROUND

101 2.1 Magnolia

102 The phrase *generic programming* has over decades of programming language development come to
103 have a variety of interpretations, depending on the main type of genericity considered. Gibbons gives a
104 taxonomy of interpretations (Gibbons, 2006). Stepanov-style generic programming (Dehnert and Stepanov,
105 1998) corresponds to what Gibbons calls *genericity by property*, where one describes data structures and

algorithms in terms of syntactic and semantic requirements. This is the essence of Stepanov’s and Musser’s *concepts* (Musser and Stepanov, 1988). They are the direct inspiration behind C++0x concepts (Gregor et al., 2006); the C++20 concepts are a scaled back realization of those that only allow syntactic requirements on instantiations. (In this latter case, we talk of *genericity by structure*.)

Magnolia is a programming language designed as an embodiment of Stepanov-style generic programming (Bagge, 2009). Magnolia code is structured into modules that mix abstract specifications of operations and their concrete implementations flexibly, following the work of Goguen and Burstall on the theory of institutions (Goguen and Burstall, 1984). The language does not offer any primitive types aside from predicates: every data structure is implemented in a configurable host programming language. As of today, Magnolia can target C++ and Python (Chetioui, 2021). Our prior work coins the term *genericity by host language* to refer to this axis of parameterization, in the style of Gibbons’ taxonomy (Chetioui et al., 2022). Composite operations can be implemented in Magnolia, while the base types and operations, including loop structures, are implemented in the host language. The programmer can freely decide where to set the boundary between the operations implemented in Magnolia, and those implemented in the base library written in the host language — depending on what is more convenient. An appropriate choice of underlying data structures results in code that is as performant as if implemented directly in the host language (Chetioui et al., 2022). Because the axiom formalism is semantically compatible with the program code, Magnolia avoids the semantic gap common in approaches to formal software verification (Sannella and Tarlecki, 1996).

A Magnolia **signature** declares types and operations. A **signature** can be augmented with **axioms** that restrict the properties of its types and operations: the resulting module is a **concept**. An **implementation** allows the same declarations as a **signature**, but also (generic) implementations for the declared operations. The last kind of module in Magnolia is a **program**, a specific kind of **implementation** in which all the specified operations and types are matched with implementations. Crucially, types and operations in a **program** are no longer generic but instead fully concrete. An **implementation** can be a model of a **concept**; a **concept** can also be a model of another **concept**. Such modeling relations can be specified directly in Magnolia using the **satisfaction** language construct.

Magnolia operations can be **functions**, **procedures**, and **predicates**. The arguments to functions and predicates are immutable, while arguments to procedures are given explicit modes: **obs** (read-only), **upd** (read/write), and **out** (write-only, and the procedure promises to initialize the argument). Procedures do not return a value. Calls to procedures are prefixed with the **call** keyword.

Listing 1 gives a general overview of the different kinds of Magnolia modules. We first specify the signature of a magma (a set T with a closed binary operation `bop`). By asserting the associativity property on a magma, we get a semigroup. The `ConcretePartialSemigroup` implementation describes an external C++ API providing a guarded multiplication operator over integer matrices, where the guard is intended to ensure the argument matrices have compatible dimensions. `ExampleProgram` builds `multiplyThreeMatrices` off of the primitive building blocks provided by `ConcretePartialSemigroup`. The `ExampleProgramHasMulPartialSemigroup` **satisfaction** relation indicates that `ExampleProgram` satisfies the semigroup axioms, with the set of integer matrices and guarded multiplication on it. The guard provided on the multiplication operation in the left-hand side of the satisfaction is propagated to the right-hand side. The resulting satisfaction relation asserts the `ExampleProgram` has a partial semigroup structure. A block of renamings (`[T => IntMatrix, bop => *_]`) is applied to `Semigroup`. Magnolia’s renamings allow changing the names of types and operations in places where a module is “opened”. This is a powerful feature which allows normalizing

the names exposed by modules when we open them in a given scope, independently of how their types and operations were initially named.

Listing 1. Multiplying three matrices in Magnolia.

```

152 signature Magma = {
153   type T;
154   function bop(a: T, b: T): T;
155 }
156
157 concept Semigroup = {
158   use Magma;
159   axiom associativity(a: T, b: T, c: T) {
160     assert bop(bop(a, b), c) == bop(a, bop(b, c));
161   }
162 }
163
164 implementation ConcretePartialSemigroup =
165   external C++ lib.int_matrices {
166     type Nat;
167     type IntMatrix;
168
169     predicate lhsNrowsIsRhsNcols(m1: IntMatrix, m2: IntMatrix);
170
171     function _*_ (m1: IntMatrix, m2: IntMatrix): IntMatrix
172       guard lhsNrowsIsRhsNcols(m1, m2);
173   }
174
175 program ExampleProgram = {
176   use ConcretePartialSemigroup;
177
178   function multiplyThreeMatrices(
179     A: IntMatrix, B: IntMatrix, C: IntMatrix): IntMatrix = A * B * C;
180 }
181
182 satisfaction ExampleProgramHasMulPartialSemigroup = ExampleProgram
183   models Semigroup[ T => IntMatrix, bop => _*_ ];

```

2.1.1 Exploiting Magnolia axioms

Concept axioms have previously found use as test oracles (Bagge et al., 2011) and as generic optimization rules (Tang and Järvi, 2015; Bagge and Haverlaen, 2009). We implement two module transformations called *rewrite* and *generate* in the Magnolia compiler under active development (Chetioui, 2021).

The *rewrite* transformation extracts all assertions of equations from a given concept, and uses them as directed rewrite rules within a target module expression. The maximum allowed number of applications of these directed rewrite rules is provided as an argument to the transformation.

191 The *generate* transformation highlights a third possible use case for Magnolia axioms, i.e. code generation.
 192 The transformation extracts all the assertions of equations from a given concept where the left-hand side
 193 is a call to a declared function (or predicate) with two-by-two distinct universally quantified arguments,
 194 and generates an implementation for the function where the body is the right-hand side of the assertion.
 195 Intuitively, an assertion with the properties we outlined describes the behavior of the function on the
 196 left-hand side at every point. Therefore, such assertions are not only a way to specify the intended behavior
 197 of a function, but also a way to derive an actual implementation for it in case one was not already provided.
 198 Figure 1 describes the grammar for the *rewrite* and *generate* transformations.

Figure 1. The grammar for the *rewrite* and *generate* module transformations in Magnolia.

```

<transformation> ::= 'rewrite' <module-expr> 'with' <module-expr> <int>
| 'generate' <module-expr> 'in' <module-expr>

```

199 Consider the `multiplyThreeMatrices` function in Listing 1. The function is intended to multiply
 200 three matrices together — and its body `A * B * C` desugars to the expression `__*__(__*__(A, B), C)`.
 201 Due to the associativity property, the order in which the multiplications are executed does not matter when
 202 it comes to the correctness of the result. However, it matters a lot when it comes to performance: suppose A
 203 is of dimensions 100×2 , B of dimensions 2×20 , and C of dimensions 20×90 . Executing `A * B` requires
 204 $100 \times 2 \times 20$ scalar multiplications, and executing `(A * B) * C` thus requires $100 \times 2 \times 20 + 100 \times$
 205 $20 \times 90 = 184000$ scalar multiplications. On the other side, executing `B * C` requires $2 \times 20 \times 90$ scalar
 206 multiplications, and executing `A * (B * C)` requires executing $2 \times 20 \times 90 + 100 \times 2 \times 90 = 21600$
 207 scalar multiplications, nearly ten times fewer.

208 Suppose that a developer wants to use the `multiplyThreeMatrices` function in their program.
 209 They care about efficiency, and know that the input matrices A, B, and C have the same dimensions as
 210 specified above. They can use the assertion provided in the associativity property of the `Semigroup`
 211 concept as a rewrite rule in `multiplyThreeMatrices` to optimize the expression from `(A * B) * C`
 212 to `A * (B * C)`. Listing 2 shows how.

Listing 2. Demonstration of the Magnolia *rewrite* transformation.

```

213 program DevProgram = rewrite ExampleProgram
214   with Semigroup[ bop => __*__, T => IntMatrix ] 1;

```

215 The Magnolia *rewrite* module transformation takes three arguments: the module on which to perform
 216 the rewrite (`ExampleProgram` in the example), the module from which to extract rewriting rules
 217 (`Semigroup` with some renamings applied in the example), and a maximum allowed number of rule
 218 applications (1 in the example).

219 Here, `multiplyThreeMatrices` is a toy example, and defined directly in the **program** being
 220 transpiled — it would therefore be very easy to reimplement it manually. However, this is not always the
 221 case: the function one wants to transform could be very complicated, and hidden deep inside an external
 222 dependency. Without the ability to perform rewritings on functions that have been previously defined, the
 223 developer would have to write their own version of this function.

3 METHODOLOGY AND CASE STUDY

We describe here our proposed methodology for writing performant and portable code productively using the Magnolia programming language. Each step is first described from a high-level perspective, and then concretely demonstrated for our PDE solver example.

3.1 Identifying and Formalizing the Domain

The first step of our methodology is to build a thorough understanding of the targeted problem. We do that by identifying and formalizing the domain underlying the problem. Formalizing the domain gives us a mathematical understanding of the properties expected of the types and operations involved in the problem. These in turn allow specifying semantics-preserving optimization rules on them, whose correctness can be proven.

PDE solvers using FDM are based on multi-dimensional array computations. In 2018, Burrows et al. identified an array API for FDM solvers. In 2019, Chetioui et al. followed up with a formalization of the identified array API using MoA. We will first give an overview of PDE solvers as described by Burrows et al., and an introduction to the corresponding MoA theory. With this background in place, we will reimplement the PDE solver based on FDM from the work of Chetioui et al., and implement hardware-agnostic and hardware-dependent rewriting rules. We show how they can be applied to our Magnolia program, and measure the resulting performance improvements.

3.1.1 PDEs

PDE solvers have many application areas. One example is numerical simulations of wind flow — e.g. for optimizing windmill positioning in large-scale wind farms.

Computing solutions to PDEs numerically requires discretizing continuous equations to a discrete domain. This approach to PDE solvers is often illustrated in the literature with Burgers' equation (Burgers, 1948). Equation 1 presents the equation in its coordinate-free form.

$$\frac{\partial \vec{u}}{\partial t} + \vec{u} \cdot \nabla \vec{u} = \nu \nabla^2 \vec{u}, \quad (1)$$

where \vec{u} is velocity, t time, and ν the viscosity coefficient.

Assuming a 3D space, we can use a Cartesian coordinate system to rewrite Equation 1 as the following system of equations

$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} + v \frac{\partial u}{\partial y} + w \frac{\partial u}{\partial z} = \nu \frac{\partial^2 u}{\partial x^2} + \nu \frac{\partial^2 u}{\partial y^2} + \nu \frac{\partial^2 u}{\partial z^2} \quad (2)$$

$$\frac{\partial v}{\partial t} + u \frac{\partial v}{\partial x} + v \frac{\partial v}{\partial y} + w \frac{\partial v}{\partial z} = \nu \frac{\partial^2 v}{\partial x^2} + \nu \frac{\partial^2 v}{\partial y^2} + \nu \frac{\partial^2 v}{\partial z^2} \quad (3)$$

$$\frac{\partial w}{\partial t} + u \frac{\partial w}{\partial x} + v \frac{\partial w}{\partial y} + w \frac{\partial w}{\partial z} = \nu \frac{\partial^2 w}{\partial x^2} + \nu \frac{\partial^2 w}{\partial y^2} + \nu \frac{\partial^2 w}{\partial z^2}, \quad (4)$$

where $\vec{u} = (u, v, w)$.

252 To discretize the domain, we describe a $N_x \times N_y \times N_z$ grid of velocity values bounded by L_x (respectively
 253 L_y and L_z) on axis x (respectively y and z) such that the u component of the velocity at index (i, j, k) and
 254 timestep n is given by

$$u_{i,j,k}^n = u(i\Delta x, j\Delta y, k\Delta z, n\Delta t), \quad (5)$$

255 with $\Delta x = \frac{L_x}{N_x}$, $\Delta y = \frac{L_y}{N_y}$, and $\Delta z = \frac{L_z}{N_z}$.

256 Similarly, the partial derivative of u in the x direction at index (i, j, k) and timestep $n + 1$ is

$$\frac{\partial u}{\partial x}(i\Delta x, j\Delta y, k\Delta z, (n + 1)\Delta t). \quad (6)$$

257 In the FDM, we compute a partial derivative as a weighted sum of neighbouring grid points — where
 258 the weights are given by a list of factors called a *stencil*. The stencil is chosen by a numerical expert. This
 259 paper, following the work of Burrows et al. uses the numerical stencils $(-\frac{1}{2}, 0, \frac{1}{2})$ and $(1, -2, 1)$ for the
 260 first and second order partial derivatives respectively.

261 Given these stencils, the partial derivative of u in the x direction at index (i, j, k) and timestep $n + 1$ is
 262 approximated by

$$\frac{\partial u}{\partial x}(i\Delta x, j\Delta y, k\Delta z, (n + 1)\Delta t) \approx \frac{\Delta t}{2\Delta x}(u_{i+1,j,k}^n - u_{i-1,j,k}^n), \quad (7)$$

263 which is accurate to $O((\Delta x)^2, \Delta t)$. Computing the partial derivative along the y (respectively z) axis
 264 follows a similar pattern, where j (respectively k) varies instead of i .

265 The standard 3D explicit finite difference approximation of Equation 2 is then given by

$$\begin{aligned} u_{i,j,k}^{n+1} = & u_{i,j,k}^n - \frac{\Delta t}{2\Delta x} u_{i,j,k}^n (u_{i+1,j,k}^n - u_{i-1,j,k}^n) + \frac{\nu \Delta t}{(\Delta x)^2} (u_{i+1,j,k}^n + u_{i-1,j,k}^n - 2u_{i,j,k}^n) \\ & - \frac{\Delta t}{2\Delta y} v_{i,j,k}^n (u_{i,j+1,k}^n - u_{i,j-1,k}^n) + \frac{\nu \Delta t}{(\Delta y)^2} (u_{i,j+1,k}^n + u_{i,j-1,k}^n - 2u_{i,j,k}^n) \\ & - \frac{\Delta t}{2\Delta z} w_{i,j,k}^n (u_{i,j,k+1}^n - u_{i,j,k-1}^n) + \frac{\nu \Delta t}{(\Delta z)^2} (u_{i,j,k+1}^n + u_{i,j,k-1}^n - 2u_{i,j,k}^n). \end{aligned}$$

266 The discretization of Equations 3 and 4 follows the same pattern.

267 The API of Burrows et al. is sufficient to compute numerical solutions to PDEs using FDM. It consists
 268 of elementwise arithmetic operations at the array level $(+, -, *)$, a rotation operation on arrays (called
 269 “shift”), and arithmetic operations at the scalar level — corresponding to the behavior of the elementwise
 270 operations at each index of the array.

271 3.1.2 MoA

272 MoA (Mullin, 1988; Mullin and Jenkins, 1996) is an algebra for describing operations on arrays. MoA
 273 distinguishes between two abstraction levels: the *Denotational Normal Form* (DNF), which describes an
 274 array by its shape together with a function describing its value at every index, and the *Operational Form*
 275 (OF) which describes it on the level of the memory layout. Programs written at the DNF level do not
 276 presume knowledge of a hardware architecture. Reasoning at the DNF level is thus completely hardware

agnostic. By repeatedly applying a set of terminating rewrite rules, any array expression can be reduced to its DNF (Mullin and Thibault, 1994; Chetioui et al., 2019) — where the resulting array is described at each index by indexing into the input arrays and scalar-level operations.

Given information about the hardware architecture and the memory layout of the arrays, the ψ -correspondence theorem (Mullin and Jenkins, 1996) allows transforming a DNF expression into a corresponding hardware-dependent OF — in which the access patterns on the array are described in terms of *start*, *stride*, and *length*.

Chetioui et al. investigate the fragment of MoA corresponding to the API identified by Burrows et al., and show that for programs based on it, DNF reduction is indeed canonical, which draws appeal to MoA as a framework for the optimization of PDE solvers based on FDM.

We give an informal overview of some operations at the DNF and OF levels below. We refer the interested reader to previous work for formal definitions (Chetioui et al., 2021; Mullin, 1988).

DNF Operations

The *dimension* of an array A is denoted $\dim(A)$. It corresponds to the number of axes of the array. For $\dim(A) = n$, the *shape* of A is an n -element vector $\rho(A) = \langle s_0, \dots, s_{n-1} \rangle$ where s_i is the length of axis i . The total number of elements (or *size*) of A is given by the product of the shape, $\Pi\rho(A) = \prod_{i=0}^{n-1} s_i$.

In the definitions below A stands for an arbitrary array with dimension n and shape as defined above. Further, we use the following array in examples:

$$M = \begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix}$$

Thus, $\rho(M) = \langle 3, 2 \rangle$.

The relevant MoA operations on the DNF level are:

- the indexing function ψ , which takes an index i into an array and returns the subarray at the indexed position. When i 's length equals the dimension of the array, i is a *total* index. Otherwise, it is *partial*. $\langle \rangle \psi A = A$ holds. For our example, we have

$$\langle 2 \rangle \psi M = \langle 5, 6 \rangle$$

$$\rho(\langle 2 \rangle \psi M) = \langle 2 \rangle$$

- the reshape function that takes an array A and a shape s such that $\Pi s = \Pi\rho(A)$, and creates a new array with shape s containing the elements of A . Thus, $\rho(\text{reshape}(s, A)) = s$ holds. For example,

$$\text{reshape}(M, \langle 2, 3 \rangle) = \begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix}$$

- a rotation function rotate that takes an array A , an axis j and an offset o , and shift A by o along its j^{th} axis. The shape is unchanged, i.e. $\rho(\text{rotate}(A, j, o)) = \rho(A)$ holds. We give a few examples of how

rotation behaves on axis 0 and 1 of M:

$$\text{rotate}(M, 0, 1) = \begin{pmatrix} 5 & 6 \\ 1 & 2 \\ 3 & 4 \end{pmatrix},$$

$$\text{rotate}(M, 0, -1) = \begin{pmatrix} 3 & 4 \\ 5 & 6 \\ 1 & 2 \end{pmatrix},$$

$$\text{rotate}(M, 1, 1) = \begin{pmatrix} 2 & 1 \\ 4 & 3 \\ 6 & 5 \end{pmatrix}.$$

297 ψ -Reduction

298 Mullin and Thibault described a rewriting system for MoA expressions at the DNF level, referred to as
 299 ψ -reduction. They conjectured that ψ -reduction is canonical — and thus takes any expression to its unique
 300 DNF. This conjecture was proven to hold by Chetioui et al. for the fragment of the rewriting system required
 301 by the array API identified by Burrows et al. (Chetioui et al., 2019). ψ -reduction essentially consists of rules
 302 that move indexing operations inwards — until eventually, the expression does not contain any collective
 303 operation, but consists only of indexing and scalar operations. As a consequence, it is guaranteed that the
 304 resulting array expression can be computed without the need to materialize any intermediate array. Because
 305 the rewriting system is canonical, another consequence is that the form in which we choose to express
 306 our computation is irrelevant: all equivalent expressions in the language of MoA reduce to the same DNF
 307 expression.

308 OF Operations

309 At the OF level, we assume knowledge of the target architecture, and an intended memory layout of the
 310 array. The central MoA operations on the OF level are:

- 311 • the family of lifting operations lift_j that take two natural numbers d, q such that $d \cdot q = s_j$, and reshape
 312 A into the shape $\langle s_0, \dots, s_{j-1}, d, q, s_{j+1}, \dots, s_{n-1} \rangle$;
- 313 • the flattening function rav that transforms a multidimensional array into its linear representation in
 314 memory. Thus, $\rho(\text{rav}(A)) = \langle \Pi \rho(A) \rangle$ holds;
- 315 • the mapping function γ , which takes a shape s with $\Pi s = \Pi \rho(A)$ and a total index into A and returns
 316 the corresponding index into $\text{rav}(A)$. In this paper, we assume a row-major ordering.

317 The OF operations presented here are crucial to the theory of MoA. We thus include them for the sake of
 318 completion. These operations however do not appear explicitly in the development of our methodology.

319 3.1.3 Initial Magnolia Implementation

320 We implemented a PDE solver using the MoA array API. The implementation consists of four
 321 components:

- 322 1. a specification of the necessary MoA types and operations, with axioms asserting that they respect the
 323 relevant properties;
- 324 2. a foreign API exposing the core types and operations of the MoA specification;

- 325 3. an external implementation of the foreign API in a host language (C++);
- 326 4. an implementation of the PDE solver built upon the external MoA building blocks.

327 The ψ -calculus conflates arrays, indices, shapes, and scalars into a single array type. While convenient in
 328 the formalism, we distinguish these types in our Magnolia implementation for ease of reasoning, and to
 329 leverage the language's type system to avoid programming errors.

330 Listing 3 shows the API from Burrows et al. in the language of MoA.

Listing 3. An array API for FDM solvers in Magnolia.

```

331 signature ArrayAPI = {
332   type Array;
333   type E;
334
335   type Axis;
336   type Index;
337   type Offset;
338
339   /* Scalar-Scalar operations */
340   function _+_ (lhs: E, rhs: E): E;
341   function _- (lhs: E, rhs: E): E;
342   function _* (lhs: E, rhs: E): E;
343   function _/ (lhs: E, rhs: E): E;
344
345   /* Scalar-Array operations */
346   function _+_ (lhs: E, rhs: Array): Array;
347   // ... prototypes as above for _-, _*, _/
348
349   /* Array-Array operations */
350   function _+_ (lhs: Array, rhs: Array): Array;
351   // ... prototypes as above for _-, _*, _/
352
353   /* Rotation */
354   function rotate(array: Array, axis: Axis, offset: Offset): Array;
355
356   /* Indexing */
357   function psi(ix: Index, array: Array): E;
358 }
```

359 The declaration of the types and operations form an algebraic *signature*. We augment that signature with
 360 semantic properties in the form of *axioms* to obtain a *concept*. Listing 4 relates each array-level arithmetic
 361 operation in the API to its corresponding scalar-level operation (Burrows et al., 2018; Chetioui et al., 2019).
 362 The axioms for all binary operations follow the same pattern, we hence only show axiom bodies for the +
 363 operation for the sake of brevity.

Listing 4. Axioms for the arithmetic operations of our array API.

```

364 concept ArrayAPI_ArithmeticAxioms = {
365   require ArrayAPI;
```

```

366
367  /* Scalar-Array Axioms */
368  axiom scalarBinaryPlusAxiom(lhs: E, rhs: Array, ix: Index) {
369      assert psi(ix, lhs + rhs) == lhs + psi(ix, rhs);
370  }
371  // axiom scalarBinarySubAxiom(lhs: E, rhs: Array, ix: Index)
372  // axiom scalarMulAxiom(lhs: E, rhs: Array, ix: Index)
373  // axiom scalarDivAxiom(lhs: E, rhs: Array, ix: Index)
374
375  /* Array-Array Axioms */
376  axiom arrayBinaryPlusAxiom(lhs: Array, rhs: Array, ix: Index) {
377      assert psi(ix, lhs + rhs) == lhs + psi(ix, rhs);
378  }
379  // axiom arrayBinarySubAxiom(lhs: Array, rhs: Array, ix: Index)
380  // axiom arrayMulAxiom(lhs: Array, rhs: Array, ix: Index)
381  // axiom arrayDivAxiom(lhs: Array, rhs: Array, ix: Index)
382  }

```

383 The specifications in Listing 3 are (straightforwardly) implemented as external C++ functions and types,
 384 not shown here. Lastly, Listing 5 shows our implementation of one full step of the PDE.

Listing 5. Implementation of one full step of the PDE solver in Magnolia.

```

385  /* Solver */
386  procedure step(upd u0: Array, upd u1: Array, upd u2: Array,
387               obs nu: Float, obs dx: Float, obs dt: Float) {
388      var _1 = one(): Float;
389      var _2 = two(): Float;
390
391      var c0 = _1/_2/dx;
392      var c1 = _1/dx/dx;
393      var c2 = _2/dx/dx;
394      var c3 = nu;
395      var c4 = dt/_2;
396
397      call allSubsteps(u0, u1, u2, c0, c1, c2, c3, c4);
398  }
399
400  procedure allSubsteps(upd u0: Array, upd u1: Array, upd u2: Array,
401                      obs c0: Float, obs c1: Float, obs c2: Float,
402                      obs c3: Float, obs c4: Float) {
403      var v0 = u0;
404      var v1 = u1;
405      var v2 = u2;
406
407      v0 = substep(v0, u0, u0, u1, u2, c0, c1, c2, c3, c4);
408      v1 = substep(v1, u1, u0, u1, u2, c0, c1, c2, c3, c4);

```



```

409     v2 = substep(v2, u2, u0, u1, u2, c0, c1, c2, c3, c4);
410     u0 = substep(u0, v0, u0, u1, u2, c0, c1, c2, c3, c4);
411     u1 = substep(u1, v1, u0, u1, u2, c0, c1, c2, c3, c4);
412     u2 = substep(u2, v2, u0, u1, u2, c0, c1, c2, c3, c4);
413 }
414
415 function substep(u: Array, v: Array, u0: Array,
416                 u1: Array, u2: Array, c0: Float,
417                 c1: Float, c2: Float, c3: Float,
418                 c4: Float): Array =
419     u + c4 * (c3 * (c1 *
420         (rotate(v, zero(), -one(): Offset) +
421         rotate(v, zero(), one(): Offset) +
422         rotate(v, one(): Axis, -one(): Offset) +
423         rotate(v, one(): Axis, one(): Offset) +
424         rotate(v, two(): Axis, -one(): Offset) +
425         rotate(v, two(): Axis, one(): Offset)) - three() * c2 * u0) -
426     c0 * ((rotate(v, zero(), one(): Offset) -
427         rotate(v, zero(), -one(): Offset)) * u0 +
428         (rotate(v, one(): Axis, one(): Offset) -
429         rotate(v, one(): Axis, -one(): Offset)) * u1 +
430         (rotate(v, two(): Axis, one(): Offset) -
431         rotate(v, two(): Axis, -one(): Offset)) * u2));
432
433 /* Float utils */
434 require function _(f: Float): Float;
435 require function one(): Float;
436 require function two(): Float;
437 require function three(): Float;
438
439 /* Axis utils */
440 require function zero(): Axis;
441 require function one(): Axis;
442 require function two(): Axis;
443
444 /* Offset utils */
445 require function one(): Offset;
446 require function _(o: Offset): Offset;

```

447 3.2 Deriving Optimization Rules

448 Armed with a thorough understanding of the problem, we can now derive semantics-preserving
449 optimization rules — hardware-specific or otherwise.

Before we can apply rewriting rules defined using MoA to our program, we need to change its level of abstraction, i.e. go from an implementation that describes the resulting array using whole-array operations to one that describes its value at every index. Consider the `ToIwiseGenerator` concept in Listing 6.

The `toIwiseGenerator` axiom consists of a single assertion, which describes the behavior of the `substepIx` function when all of its arguments are universally quantified distinct variables. The right-hand side of the equation is thus a valid implementation for `substepIx`. Because this function is not implemented in the original program, we can use the *generate* transformation with `ToIwiseGenerator` to generate an implementation of `substepIx` in the implementation given in Listing 5. So as to enable further optimizations, *generate* unfolds function calls in the right-hand side of the equation. The resulting index-level code is shown in Listing 7.

Listing 6. A generator for an index-level implementation of `substep`.

```

460 concept ToIwiseGenerator = {
461     type Array;
462     type Float;
463     type Index;
464
465     function substepIx(u: Array, v: Array, u0: Array,
466                     u1: Array, u2: Array, c0: Float,
467                     c1: Float, c2: Float, c3: Float,
468                     c4: Float, ix: Index): Float;
469
470     function substep(u: Array, v: Array, u0: Array,
471                    u1: Array, u2: Array, c0: Float,
472                    c1: Float, c2: Float, c3: Float,
473                    c4: Float): Array;
474
475     function psi(ix: Index, array: Array): Float;
476
477     axiom toIwiseGenerator(u: Array, v: Array, u0: Array,
478                          u1: Array, u2: Array, c0: Float,
479                          c1: Float, c2: Float, c3: Float,
480                          c4: Float, ix: Index) {
481         assert substepIx(u, v, u0, u1, u2, c0, c1, c2, c3, c4, ix) ==
482             psi(ix, substep(u, v, u0, u1, u2, c0, c1, c2, c3, c4));
483     }
484 }
```

Listing 7. Index-level implementation of `substep` in Magnolia.

```

485
486 require function psi(ix: Index, array: Array): Float;
487
488 function substepIx(u: Array, v: Array, u0: Array,
489                  u1: Array, u2: Array, c0: Float,
490                  c1: Float, c2: Float, c3: Float,
491                  c4: Float, ix: Index): Float =
```

```

492   psi(ix, u + c4 * (c3 * (c1 *
493       (rotate(v, zero(): Axis, -one(): Offset) +
494       rotate(v, zero(): Axis, one(): Offset) +
495       rotate(v, one(): Axis, -one(): Offset) +
496       rotate(v, one(): Axis, one(): Offset) +
497       rotate(v, two(): Axis, -one(): Offset) +
498       rotate(v, two(): Axis, one(): Offset)) -
499   three() * c2 * u0) - c0 *
500       ((rotate(v, zero(): Axis, one(): Offset) -
501       rotate(v, zero(): Axis, -one(): Offset)) * u0 +
502       (rotate(v, one(): Axis, one(): Offset) -
503       rotate(v, one(): Axis, -one(): Offset)) * u1 +
504       (rotate(v, two(): Axis, one(): Offset) -
505       rotate(v, two(): Axis, -one(): Offset)) * u2)))));
506 }

```

507 To make use of `substepIx` within the program, we need to replace calls to `substep` with calls to a
508 scheduling function that uses `substepIx` to describe the value of the array at every index. We use the
509 program transformation **rewrite** ... **with** `ToIxwise 1` to achieve that, with `ToIxwise` a concept
510 of Listing 8. Throughout the rest of the paper, we use the term *schedule* like in Halide (Ragan-Kelley et al.,
511 2012).

Listing 8. A concept with a rewrite rule from `substep` to a new scheduling function.

```

512 concept ToIxwise = {
513     type Array;
514     type Float;
515
516     function substep(u: Array, v: Array, u0: Array,
517         u1: Array, u2: Array, c0: Float,
518         c1: Float, c2: Float, c3: Float,
519         c4: Float): Array;
520
521     function schedule(u: Array, v: Array, u0: Array,
522         u1: Array, u2: Array, c0: Float,
523         c1: Float, c2: Float, c3: Float,
524         c4: Float): Array;
525
526     axiom toIxwiseRule(u: Array, v: Array, u0: Array,
527         u1: Array, u2: Array, c0: Float,
528         c1: Float, c2: Float, c3: Float,
529         c4: Float) {
530         assert substep(u, v, u0, u1, u2, c0, c1, c2, c3, c4) ==
531             schedule(u, v, u0, u1, u2, c0, c1, c2, c3, c4);
532     }
533 }

```

534 Magnolia does not expose native looping constructs. For that reason, the implementation of `schedule`
 535 is done in the host language. The `schedule` function uses the `substepIx` function in Listing 7 to
 536 describe the content of the result array at every index.

537 From that point onwards, we can use MoA to derive transformation rules on our program.

538 3.2.1 Hardware-Agnostic Transformation Rules

539 In their work on embedding Burrows et al.’s array API for FDM solvers in MoA, Chetioui et al. outline a
 540 rewriting system sufficient to transform a program based on this API to its DNF. This rewriting system is
 541 canonical, i.e. rewriting always terminates, and the order in which the rules are applied is inconsequential.

542 Rewriting rules at the DNF level do not require hardware knowledge, and therefore constitute hardware-
 543 agnostic transformation rules. We show an implementation of these rules in Magnolia in Listing 9.

Listing 9. The DNF rewriting rules in Magnolia.

```

544 concept GenericBinopRules = {
545   type E;
546   type Array;
547   type Index;
548
549   function binop(lhs: E, rhs: E): E;
550   function binop(lhs: E, rhs: Array): Array;
551   function binop(lhs: Array, rhs: Array): Array;
552   function psi(ix: Index, array: Array): E;
553
554   // Rule 1
555   axiom binopArrayRule(ix: Index, lhs: Array, rhs: Array) {
556     assert psi(ix, binop(lhs, rhs)) ==
557       binop(psi(ix, lhs), psi(ix, rhs));
558   }
559
560   // Rule 2
561   axiom binopScalarRule(ix: Index, lhs: E, rhs: Array) {
562     assert psi(ix, binop(lhs, rhs)) == binop(lhs, psi(ix, rhs));
563   }
564 }
565
566 concept DNFRules = {
567   use GenericBinopRules[ binop => _+_
568                       , binopScalarRule => addScalarRule
569                       , binopArrayRule => addArrayRule
570                       ];
571   use GenericBinopRules[ binop => _-_
572                       , binopScalarRule => subScalarRule
573                       , binopArrayRule => subArrayRule
574                       ];
575   use GenericBinopRules[ binop => _*_

```

```

576         , binopScalarRule => mulScalarRule
577         , binopArrayRule => mulArrayRule
578     ];
579     use GenericBinopRules[ binop => _/_
580         , binopScalarRule => divScalarRule
581         , binopArrayRule => divArrayRule
582     ];
583
584     type Axis;
585     type Offset;
586
587     function rotate(array: Array, axis: Axis, offset: Offset): Array;
588     function rotateIx(ix: Index, axis: Axis, offset: Offset): Index;
589
590     // Rule 3
591     axiom rotateRule(ix: Index, array: Array, axis: Axis,
592         offset: Offset) {
593         assert psi(ix, rotate(array, axis, offset)) ==
594             psi(rotateIx(ix, axis, offset), array);
595     }
596 }[ E => Float ];

```

As explained in Subsubsection 3.1.2, applying the DNF rules pushes computations down from the array-level to the index-level, i.e. the resulting computations are devoid of whole-array operations and contain only indexing and scalar arithmetic operations.

Table 1 shows runtime results for our PDE solver implementation in Magnolia, before and after full DNF reduction using the DNF rewriting rules. DNF reduction speeds up the code by a factor of roughly $4.18\times$ and significantly reduces memory usage. At the DNF level, the expression is written in terms of scalar and indexing operations, eliminating the need to compute temporary arrays, and increasing computational density. This experiment shows that such a rewriting system gives the ability to write programs using whole-array operations without losing out on the benefits of writing index-level code. The ability to write algorithms in different ways without inducing a loss of performance is key to the productive development of performant code.

	Wall time (in seconds)
Before DNF reduction	323.02
After DNF reduction	42.26

Table 1. Execution time (in seconds) of the 3-dimensional PDE solver Magnolia implementation compiled to C++, with and without reduction to DNF. The code is compiled with gcc 10.2.1 with optimization level O3. The space dimensions are $256 \times 256 \times 256$ and the solver is run for 50 timesteps. The code is run 10 times on the Intel Xeon Silver 4112 CPU, and the time measurements are averaged.

3.2.2 Hardware-Specific Transformation Rules

Which hardware-specific transformation rules are relevant to implement is by nature dependent on the underlying hardware architecture we are interested in. For example, Chetioui et al.'s previous work on

611 formalizing PDE computations in MoA gave rise to rules for introducing padding into array expressions.
 612 Their work also discusses rewrites rules that use the *dimension lifting* operation, which is a *reshape*
 613 operation with the explicit purpose of matching the shape of arrays with characteristics of the underlying
 614 hardware. E.g. lifting by d_1 across the first axis allows one to *scatter* the resulting subarrays across
 615 d_1 processes; or, lifting by 4 across the last axis of an array of 32-bit floats allows one to vectorize
 616 computations on an architecture with 128-bit vector registers. The hardware architecture combined with
 617 the data dependencies of the algorithm determine the shape and layout of the arrays.

618 We discuss two examples of such hardware-dependent rewriting systems below.

619 **Example: Dimension lifting over several cores**

620 At the DNF level, our concern was to express our algorithm in a canonical form, without paying any
 621 mind to hardware-related concerns. A contrario, our concern at the OF level is to make the best use of the
 622 hardware available. The rewrites we express are thus often concerned with changing the schedule of our
 623 computations. Scheduling is handled outside of Magnolia in our example, by the `schedule` function.

624 Listing 10 showcases a rewriting rule for moving from our initial scheduling function to one that
 625 parallelizes the computation over several cores, the number of which can be parameterized externally.

Listing 10. The rewriting rules for distributing the computation on several cores.

```

626 concept OFLiftCores = {
627   type Array;
628   type Float;
629   type Axis;
630   type Nat;
631
632   function nbCores(): Nat;
633
634   function scheduleThreaded(
635       u: Array, v: Array,
636       u0: Array, u1: Array, u2: Array,
637       c0: Float, c1: Float, c2: Float, c3: Float, c4: Float,
638       nbThreads: Nat
639   ): Array;
640
641   function schedule(
642       u: Array, v: Array,
643       u0: Array, u1: Array, u2: Array,
644       c0: Float, c1: Float, c2: Float, c3: Float, c4: Float
645   ): Array;
646
647   axiom liftCoresRule(
648       u: Array, v: Array,
649       u0: Array, u1: Array, u2: Array,
650       c0: Float, c1: Float, c2: Float, c3: Float, c4: Float
651   ) {
652     assert schedule(u, v, u0, u1, u2, c0, c1, c2, c3, c4) ==

```

```

653         scheduleThreaded(u, v, u0, u1, u2, c0, c1,
654                          c2, c3, c4, nbCores());
655     }
656 }

```

657 The implementation of the new `scheduleThreaded` function must also be provided externally.
 658 Because the schedule is separate from the algorithm, the cost of expressing scheduling rewrites is mostly
 659 the cost of implementing a new schedule. Once a schedule is implemented, it can be reused for algorithms
 660 exhibiting similar data dependency patterns, and to target similar hardware. The cost of implementing
 661 scheduling rewrites thus decreases as more schedules are implemented, and more problems are explored.

662 **Example: Padding computations**

663 Figure 2 shows the dependency patterns for one third of a half-step of the PDE across the last axis
 664 of the array. The element at index i at time $t + 1$ depends on the elements at index i , $(i - 1) \bmod N$,
 665 and $(i + 1) \bmod N$ at time t . The modulo operation serves to index the right dependencies for the first
 666 (respectively last) element of the array, where decrementing (respectively incrementing) the index would
 667 create an out-of-bounds index. Modulo operations are still expensive, even on modern hardware (Lemire
 668 et al., 2019). Additionally, if N is large, the computations at the boundary need to access elements that are
 669 far apart in memory — therefore benefitting less from data locality than the computations in the middle of
 670 the array.




Figure 2. The dependency pattern for one third of a half-step of the PDE across the last axis of the array. Each column represents an array of length N indexed from 0 to $N - 1$ for a given timestep. The element at index i of the array at time $t + 1$ depends on the elements at indices i , $(i - 1) \bmod N$ and $(i + 1) \bmod N$ of the array at time t .

671 Chetioui et al. previously showed that padding is a way to eliminate these modulo computations and to
 672 increase data locality, at the cost of duplicating data in memory (Chetioui et al., 2021).

673 Figure 3 shows the dependency patterns for one third of a half-step of the PDE across the last axis of the
 674 array when the array is padded. In that case, the computation at the boundaries of the array can be rewritten

675 to depend on three adjacent elements in the array. The modulo computation can also be eliminated. We pay
 676 for these improvements by using more space, and by refilling the padding before every timestep.

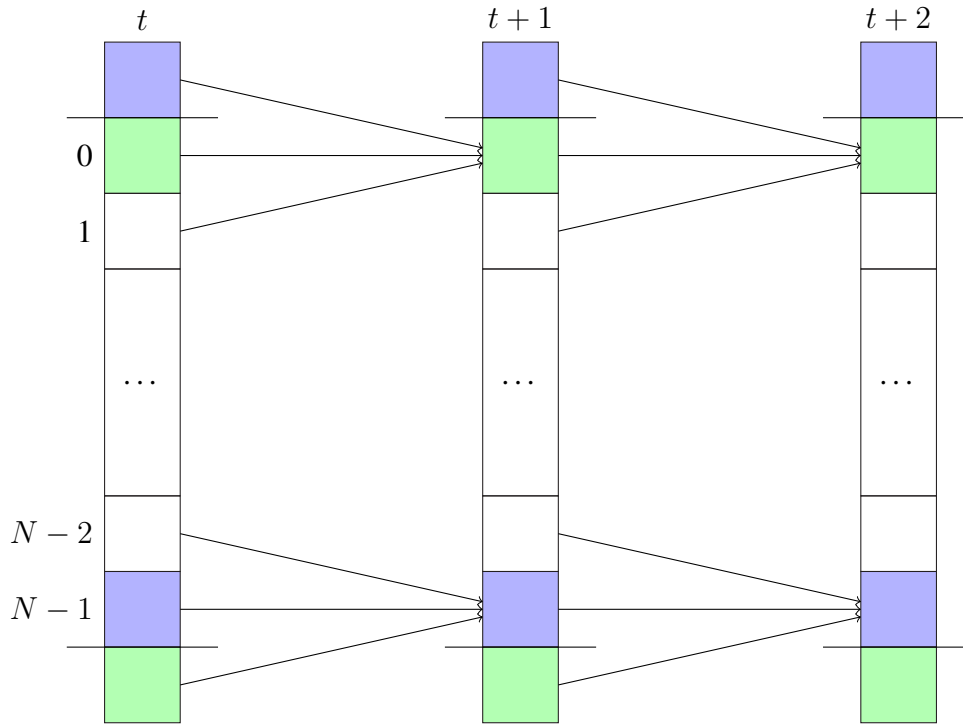


Figure 3. The dependency pattern for one third of a half-step of the PDE across the last axis of the array when the array is padded once on each side on the last axis. Each column represents an array of length N indexed from 0 to $N - 1$ for a given timestep. The elements colored in the same color have the same value. The element at index i of the array at time $t + 1$ depends on the elements at indices i , $i - 1$ and $i + 1$ of the array at time t .

677 Listing 11 shows an implementation of the padding transformation rules in Magnolia.

Listing 11. The rewriting rules for padding.

```

678 concept OFPad = {
679   type Array;
680   type Float;
681
682   procedure allSubsteps(upd u0: Array, upd u1: Array, upd u2: Array,
683                       obs c0: Float, obs c1: Float, obs c2: Float,
684                       obs c3: Float, obs c4: Float);
685
686   procedure refillPadding(upd a: Array);
687
688   function schedulePadded(u: Array, v: Array,
689                          u0: Array, u1: Array, u2: Array, c0: Float,
690                          c1: Float, c2: Float, c3: Float, c4: Float): Array;
691
692   function schedule(u: Array, v: Array,
693                    u0: Array, u1: Array, u2: Array,
```



```

694         c0: Float, c1: Float, c2: Float,
695         c3: Float, c4: Float): Array;
696
697 axiom padRule(u: Array, v: Array, u0: Array, u1: Array, u2: Array,
698             c0: Float, c1: Float, c2: Float, c3: Float,
699             c4: Float) {
700     assert schedule(u, v, u0, u1, u2, c0, c1, c2, c3, c4) ==
701     { var result =
702         schedulePadded(u, v, u0, u1, u2, c0, c1, c2, c3, c4);
703         call refillPadding(result);
704         value result;
705     };
706 }
707
708 type Index;
709 type Axis;
710 type Offset;
711 function rotateIx(ix: Index, axis: Axis, offset: Offset): Index;
712 function rotateIx_padded(ix: Index, axis: Axis, offset: Offset)
713     : Index;
714
715 axiom rotateIxPadRule(ix: Index, axis: Axis, offset: Offset) {
716     assert rotateIx(ix, axis, offset) ==
717         rotateIx_padded(ix, axis, offset);
718 }
719 }

```

720 The implementation in Listing 11 assumes that the input arrays are padded arbitrarily across each axis in
721 the host language, in a way that is compatible with the new `rotateIx_padded` function. Details such
722 as the amount of padding across each axis are therefore not visible in Magnolia. This is however purely a
723 design choice, insofar as we have chosen to make the `Index` type completely opaque. This has the benefit
724 of making the program naturally shape polymorphic to a degree — though the program is not as interesting
725 for input arrays with initial number of dimensions different than three.

726 We can control padding across each axis more explicitly by specializing our code further. This can also
727 be achieved using transformation rules — we describe the steps below.

728 Listing 12 shows an axiom following the generator pattern to specialize the shape polymorphic
729 `substepIx` to three dimensions. As previously, the call to `substepIx` on the right-hand side of
730 the equation is unfolded to enable additional optimizations.

Listing 12. A generator for a 3D implementation of `substepIx`.

```

731 concept OFSpecializeSubstepGenerator = {
732     type Index;
733     type Array;
734     type Float;
735     type ScalarIndex;

```

```

736
737 function mk_ix(i: ScalarIndex, j: ScalarIndex, k: ScalarIndex)
738     : Index;
739
740 function substepIx(u: Array, v: Array, u0: Array,
741     u1: Array, u2: Array, c0: Float, c1: Float,
742     c2: Float, c3: Float, c4: Float, ix: Index): Float;
743
744 function substepIx3D(u: Array, v: Array, u0: Array,
745     u1: Array, u2: Array, c0: Float, c1: Float, c2: Float,
746     c3: Float, c4: Float, i: ScalarIndex, j: ScalarIndex,
747     k: ScalarIndex): Float;
748
749 axiom specializeSubstepRule(u: Array, v: Array, u0: Array,
750     u1: Array, u2: Array, c0: Float, c1: Float, c2: Float,
751     c3: Float, c4: Float, i: ScalarIndex, j: ScalarIndex,
752     k: ScalarIndex) {
753     assert substepIx3D(u, v, u0, u1, u2, c0, c1, c2,
754         c3, c4, i, j, k) ==
755         substepIx(u, v, u0, u1, u2, c0, c1, c2, c3, c4,
756             mk_ix(i, j, k));
757 }
758 };

```

759 Recall the original implementation of `substepIx` given in Listing 7. Every indexing operation of some
760 array `a` in the resulting implementation of `substepIx3D` is now either of the form `psi(mk_ix(i, j`
761 `, k), a)`, or of the form `psi(rotateIx(mk_ix(i, j, k), x, o), a)` for some axis `x` and
762 some offset `o`.

763 Listing 13 introduces a specialized `psi` function for 3D arrays. It does that by introducing three
764 projection functions `ix0`, `ix1`, and `ix2` on Indexes. General indexing operations of the form
765 `psi(mk_ix(i, j, k), a)` are first specialized to expressions of the form `psi(ix0(mk_ix(i,`
766 `j, k)), ix1(mk_ix(i, j, k)), ix2(mk_ix(i, j, k)), a)` by an application of
767 `specializePsiRule` — which can then be reduced to `psi(i, j, k, a)` via three applications of
768 `reduceMakeIxRule`.

Listing 13. Specializing calls to the indexing function ψ .

```

769 concept OFSpecializePsi = {
770     type Index;
771     type Array;
772     type E;
773     type ScalarIndex;
774
775     /* 3D index projection functions */
776     function ix0(ix: Index): ScalarIndex;
777     function ix1(ix: Index): ScalarIndex;
778     function ix2(ix: Index): ScalarIndex;

```

```

779
780  /* 3D index constructor */
781  function mk_ix(i: ScalarIndex, j: ScalarIndex, k: ScalarIndex)
782      : Index;
783
784  function psi(ix: Index, array: Array): E;
785  function psi(i: ScalarIndex, j: ScalarIndex, k: ScalarIndex,
786      array: Array): E;
787
788  axiom specializePsiRule(ix: Index, array: Array) {
789      assert psi(ix, array) == psi(ix0(ix), ix1(ix), ix2(ix), array);
790  }
791
792  axiom reduceMakeIxRule(i: ScalarIndex, j: ScalarIndex,
793      k: ScalarIndex) {
794      var ix = mk_ix(i, j, k);
795      assert ix0(ix) == i;
796      assert ix1(ix) == j;
797      assert ix2(ix) == k;
798  }
799 }[ E => Float ];

```

800 We also want to call our specialized version of `psi` instead of the general one in expressions now of the
801 form `psi(ix0(rx), ix1(rx), ix2(rx), a)` where `rx = rotateIx(mk_ix(i, j, k),`
802 `x, o)`. For that purpose, we can apply the rewriting rules defined in Listing 14. These rewriting rules
803 essentially unfold `rotateIx`. All the indexing operations in `substepIx3D` now use the specialized
804 form of `psi`, and the scalar indices are either constants or of the form $(i + o) \% s$, with i a scalar
805 index, o an offset, and s the length of the relevant axis of the array.

Listing 14. A transformation rules to specialize the index rotation operation.

```

806 concept OFReduceMakeIxRotate = {
807     use signature (OFSpecializePsi);
808
809     type Axis;
810     type Offset;
811
812     function zero(): Axis;
813     function one(): Axis;
814     function two(): Axis;
815
816     function rotateIx(ix: Index, axis: Axis, offset: Offset): Index;
817
818     type AxisLength;
819
820     function shape0(): AxisLength;
821     function shape1(): AxisLength;

```

```

822  function shape2(): AxisLength;
823
824  function _+_ (six: ScalarIndex, o: Offset): ScalarIndex;
825  function _%_ (six: ScalarIndex, sc: AxisLength): ScalarIndex;
826
827  axiom reduceMakeIxRotateRule(i: ScalarIndex, j: ScalarIndex,
828    k: ScalarIndex, array: Array, o: Offset) {
829    var ix = mk_ix(i, j, k);
830    var s0 = shape0();
831    var s1 = shape1();
832    var s2 = shape2();
833
834    assert ix0(rotateIx(ix, zero(), o)) == (i + o) % s0;
835    assert ix0(rotateIx(ix, one(), o)) == i;
836    assert ix0(rotateIx(ix, two(), o)) == i;
837
838    assert ix1(rotateIx(ix, zero(), o)) == j;
839    assert ix1(rotateIx(ix, one(), o)) == (j + o) % s1;
840    assert ix1(rotateIx(ix, two(), o)) == j;
841
842    assert ix2(rotateIx(ix, zero(), o)) == k;
843    assert ix2(rotateIx(ix, one(), o)) == k;
844    assert ix2(rotateIx(ix, two(), o)) == (k + o) % s2;
845  }
846 }

```

At this point, we can reintroduce padding using the rules previously defined in Listing 11, and renaming `schedulePadded` to `schedule3DPadded`, which will need to be pulled into scope from an external implementation somewhere down the line.

We decide to implement this function externally such that the array is always circularly padded at least once on each side of each axis — a decision made based on the width of the stencil. With that knowledge, we can completely eliminate the modulo operations in `substepIx3D`. Listing 15 defines the relevant rewriting rules.

Listing 15. Elimination of the modulo operations in the program.

```

854  // We suppose here that the amount of padding is sufficient across
855  // each axis for every indexing operation.
856  concept OFeliminateModuloPadding = {
857    use signature (OFReduceMakeIxRotate);
858
859    type Array;
860    type Float;
861
862    function psi(i: ScalarIndex, j: ScalarIndex, k: ScalarIndex,
863      a: Array): Float;
864

```

```

865 axiom eliminateModuloPaddingRule(i: ScalarIndex, j: ScalarIndex,
866     k: ScalarIndex, a: Array, o: Offset) {
867     var s0 = shape0();
868     var s1 = shape1();
869     var s2 = shape2();
870
871     assert psi((i + o) % s0, j, k, a) == psi(i + o, j, k, a);
872     assert psi(i, (j + o) % s1, k, a) == psi(i, j + o, k, a);
873     assert psi(i, j, (k + o) % s2, a) == psi(i, j, k + o, a);
874 }
875 }

```

Listing 16 shows how we apply the rewriting rules defined above using Magnolia's rewriting system to build a new program. Note that, as we are in the case when an implementation for `schedulePadded` is not in scope before the rules defined in `OFPad` are applied, we can replace the `rewrite` by a simple renaming — as we do in the example. To build a valid program, we also need to pull in scope external functions, such as the relevant schedules, and `psi`. These come from `ExternalNeededFunctions` in the example.

Listing 16. Putting all the rewriting rules together.

```

882 program SpecializedAndPaddedProgram = {
883     use (rewrite
884         (rewrite
885             (rewrite
886                 (generate OFSpecializeSubstepGenerator in
887                     DNFImplementation)
888                 with OFSpecializePsi 10)
889                 with OFReduceMakeIxRotate 20)
890                 with OFPad[schedulePadded =>
891                     schedule3DPadded] 1)
892                 with OFEliminateModuloPadding 10);
893 }
894
895 use ExternalNeededFunctions; // pulling in psi, schedules, etc...
896 }

```

Table 2 gives an overview of the performance improvements brought by the rewriting rules. On the specific processor considered, padding does not seem to enable any significant speedup for our original implementation. Specializing the code to our specific 3D indexing function makes the code run faster in the unpadded case, and seems to allow a significant speedup from the unpadded case to the padded case — the generated code runs nearly twice as fast in that case.

Crucially, this performance improvement did not require any reimplementations of the core algorithm. Building our core algorithm generically allows us to introduce specialized underlying types and operations, once more information is known about our input data or the underlying hardware architecture. The Magnolia term rewriting engine then allows us to introduce new operations and to replace calls to existing concrete implementations with calls to other functions with possibly different argument lists.

	Unpadded case	Padded case
Non-specialized indexing	689.46	675.12
Specialized indexing	540.3	285.55

Table 2. Execution time (in seconds) of the 3-dimensional PDE solver Magnolia implementation compiled to C++ with specialized indexing and with or without padding. The code is compiled with gcc 10.2.0 with optimization level O3. The space dimensions are $512 \times 512 \times 512$ and the solver is run for 50 timesteps. The code is run on the ThunderX2 CN9980 CPU. In the padded case, each axis is padded circularly exactly once on both ends.

907 This is another twist of generic programming: *rewrite* and *generate* allow to replace operations (or
 908 combinations of operations) in a generic module with others that have potentially different argument lists —
 909 so long as we can describe the behavior of the former at all points in terms of calls to the new operation(s).

4 DISCUSSION AND RELATED WORK

910 We presented a methodology for solving the P³ problem on existing and emerging architectures and
 911 applied it to the domain of array computations. Instead of developing one program to target n hardware
 912 architectures, we implement a single program, along with hardware-specific rewriting rules. By relating the
 913 high-level problem to a mathematical basis, we ensure that the set of optimization rules we implement is
 914 correct, and reusable for problems that can be embedded within the same formalism.

915 Magnolia gives developers the tools to write high-level, domain-specific compilers with custom
 916 optimization rules, and a custom target language. The ability to choose flexibly to which opaque building
 917 blocks a Magnolia program reduces allows the application of optimization rules at various abstraction
 918 levels, until the boundary between Magnolia and the external primitives implemented in the host language
 919 is reached. Our approach is centered around the idea of expressing generic algorithms independently from
 920 any particular schedule, i.e. independently from any hardware abstraction.

921 As we mentioned in Section 1.1, the term *schedule* as used throughout the paper originates in the
 922 work of Ragan-Kelley et al. on Halide (Ragan-Kelley et al., 2012). SPIRAL (Puschel et al., 2005) and
 923 Sequoia (Fatahalian et al., 2006) predate Halide, but make a similar distinction between an algorithm and its
 924 mapping to the underlying hardware architecture. Halide exposes a set of scheduling primitives from which
 925 developers can build their own schedules. TVM (Chen et al., 2018) follows this idea and extends Halide’s set
 926 of scheduling primitives. The set of schedules that can be expressed in such systems is necessarily limited
 927 by the set of available scheduling primitives. Extending this set requires modifications to the language
 928 and its compiler, and is thus costly. Recent work by Liu et al. shows that carefully choosing high-level
 929 rewriting rules on schedules allows optimizing tensor programs beyond what is currently possible in these
 930 languages (Liu et al., 2022). In our system, schedules are fully specified by the developer. Compared to the
 931 approach taken by Halide or TVM, the developer has full control over how their computations are executed,
 932 but incur a higher implementation cost when no scheduling algorithm exists for their particular flavor of
 933 target hardware architecture. Adding “default” scheduling primitives to Magnolia as a convenience could
 934 improve the developer experience, and is therefore a consideration for future work.

935 MLIR (Lattner et al., 2021) makes heavy use of rewrite rules through the MLIR *PatternRewrite*
 936 infrastructure (Vasilache et al., 2022). Their design is influenced by LIFT (Steuer et al., 2017, 2015),
 937 another programming language exploiting rewrite rules for high-performance array computations. In LIFT,
 938 the application of rewrite rules is automated by a stochastic search method. Hagedorn et al. extend LIFT
 939 specifically for optimizing stencil programs (Hagedorn et al., 2018). Such rewrite approaches are so far

940 limited in that they do not always deliver high enough performance for real-world use (Hagedorn et al.,
941 2020). This is in contrast to autoscheduling in Halide, which outperforms human experts on average (Adams
942 et al., 2019). Automatic scheduling techniques are key to improving solutions to the P³ problem, and are
943 thus an important topic to further explore also for rewrite rules-based optimizers.

944 Approaches to optimization based on rewrite rules, such as the one presented here, can benefit from
945 rewriting strategies, e.g. for localizing rewrites to only a particular chunk of the input program or for
946 traversing the AST in a specific order. Kirchner gives a recent survey of strategic rewriting (Kirchner, 2015).
947 Example of tools implementing such strategies include Maude (Clavel et al., 2007; Martí-Oliet et al., 2005,
948 2009) and Stratego (Visser, 2005). Hagedorn et al. introduce a functional approach to high-performance
949 code generation based on rewriting strategies (Hagedorn et al., 2020): computations are expressed in the
950 RISE programming language, and rewrite rules and strategies in the ELEVATE strategy language. Fu et al.
951 (2021) later added a type system to ELEVATE to ensure statically that rewrites are composed correctly. As
952 shown throughout the paper, our rewriting system today only provides the ability to apply sets of rewrite
953 rules a certain number of times, in sequence. Given a rule $e_1 = e_2$, the sequence $e_1; e_1$ can be rewritten
954 to $e_2; e_1$, but not directly to $e_1; e_2$. Such a transformation can be expressed today by applying the rule
955 $e_1 = e_2$ twice, and then applying the opposite rule $e_2 = e_1$ once, but this is both embarrassingly verbose and
956 inefficient. Adding rewriting strategies to Magnolia will unlock those rewrites that are not easily accessible
957 today, and thus further improve the system's code reuse capabilities. The implementation of Magnolia
958 strategies is of particular interest, and fits into our larger project of exploring module transformations
959 through the lens of *Syntactic Theory Functors* (Haveraaen and Roggenbach, 2020).

960 For future work, we also envision the implementation of an extension to the Magnolia rewriting system
961 that supports conditional rewrite rules. Conditional equations can already be expressed in Magnolia, but
962 the rewriting system is not yet able to exploit them.

963 Whether axioms constitute valid rewriting rules is verifiable by extending Magnolia with formal
964 verification tools — insofar as the relevant properties that a program must satisfy can be derived from the
965 stated axioms about its external building blocks. The properties asserted about externally implemented
966 code can however only be assumed to hold, and constitute the trusted computing base of the whole program.
967 Work on connecting verification tools with Magnolia's specification facilities is already underway, with
968 encouraging results (Hamre, 2022).

ACKNOWLEDGEMENTS

969 The research presented in this paper has benefited from the Experiment Infrastructure for Exploration of
970 Exascale Computing (eX3), which is financially supported by the Research Council of Norway under
971 contract 270053.

REFERENCES

- 972 Abts, D., Ross, J., Sparling, J., Wong-VanHaren, M., Baker, M., Hawkins, T., et al. (2020). Think fast: A
973 tensor streaming processor (TSP) for accelerating deep learning workloads. In *2020 ACM/IEEE 47th*
974 *Annual International Symposium on Computer Architecture (ISCA)*. 145–158. doi:10.1109/ISCA45697.
975 2020.00023
- 976 Adams, A., Ma, K., Anderson, L., Baghdadi, R., Li, T.-M., Gharbi, M., et al. (2019). Learning to optimize
977 halide with tree search and random programs. *ACM Trans. Graph.* 38. doi:10.1145/3306346.3322967

- 978 Bagge, A. H. (2009). *Constructs & Concepts: Language Design for Flexibility and Reliability*. Ph.D.
 979 thesis, Research School in Information and Communication Technology, Department of Informatics,
 980 University of Bergen, Norway, PB 7803, 5020 Bergen, Norway
- 981 Bagge, A. H., David, V., and Haveraaen, M. (2011). Testing with axioms in C++ 2011. *Journal of Object*
 982 *Technology* 10, 10:1–32. doi:10.5381/jot.2011.10.1.a10
- 983 Bagge, A. H. and Haveraaen, M. (2009). Axiom-based transformations: Optimisation and testing. In
 984 *Proceedings of the Eighth Workshop on Language Descriptions, Tools and Applications (LDTA 2008)*,
 985 eds. J. J. Vinju and A. Johnstone (Elsevier), vol. 238, 17–33. doi:10.1016/j.entcs.2009.09.038
- 986 Basili, V. R., Briand, L. C., and Melo, W. L. (1996). How reuse influences productivity in object-oriented
 987 systems. *Commun. ACM* 39, 104–116
- 988 Burgers, J. M. (1948). A mathematical model illustrating the theory of turbulence. In *Advances in applied*
 989 *mechanics* (Elsevier), vol. 1. 171–199
- 990 Burrows, E., Friis, H. A., and Haveraaen, M. (2018). An array API for finite difference methods. In
 991 *Proceedings of the 5th ACM SIGPLAN International Workshop on Libraries, Languages, and Compilers*
 992 *for Array Programming* (New York, NY, USA: ACM), ARRAY 2018, 59–66. doi:10.1145/3219753.
 993 3219761
- 994 Chen, T., Moreau, T., Jiang, Z., Zheng, L., Yan, E., Cowan, M., et al. (2018). TVM: An automated
 995 end-to-end optimizing compiler for deep learning. In *Proceedings of the 13th USENIX Conference on*
 996 *Operating Systems Design and Implementation* (USA: USENIX Association), OSDI’18, 579–594
- 997 [Dataset] Chetioui, B. (2021). The Magnolia Compiler. [https://github.com/magnolia-lang/](https://github.com/magnolia-lang/magnolia-lang)
 998 [magnolia-lang](https://github.com/magnolia-lang/magnolia-lang). [Online; accessed 31-January-2022]
- 999 Chetioui, B., Abusdal, O., Haveraaen, M., Järvi, J., and Mullin, L. (2021). *Padding in the Mathematics of*
 1000 *Arrays* (New York, NY, USA: Association for Computing Machinery). 15–26
- 1001 Chetioui, B., Järvi, J., and Haveraaen, M. (2022). Revisiting language support for generic programming:
 1002 the case of Magnolia. Under review
- 1003 Chetioui, B., Mullin, L., Abusdal, O., Haveraaen, M., Järvi, J., and Macià, S. (2019). Finite difference
 1004 methods fengshui: Alignment through a mathematics of arrays. In *Proceedings of the 6th ACM SIGPLAN*
 1005 *International Workshop on Libraries, Languages and Compilers for Array Programming* (New York,
 1006 NY, USA: Association for Computing Machinery), ARRAY 2019, 2–13. doi:10.1145/3315454.3329954
- 1007 Clavel, M., Durán, F., Eker, S., Lincoln, P., Martí-Oliet, N., Meseguer, J., et al. (2007). *All about Maude*
 1008 *- a High-Performance Logical Framework: How to Specify, Program and Verify Systems in Rewriting*
 1009 *Logic* (Berlin, Heidelberg: Springer-Verlag)
- 1010 Dehnert, J. C. and Stepanov, A. A. (1998). Fundamentals of generic programming. In *Selected Papers*
 1011 *from the International Seminar on Generic Programming* (Berlin, Heidelberg: Springer-Verlag), 1–11
- 1012 Fatahalian, K., Horn, D. R., Knight, T. J., Leem, L., Houston, M., Park, J. Y., et al. (2006).
 1013 Sequoia: Programming the memory hierarchy. In *Proceedings of the 2006 ACM/IEEE Conference*
 1014 *on Supercomputing* (New York, NY, USA: Association for Computing Machinery), SC ’06, 83–es.
 1015 doi:10.1145/1188455.1188543
- 1016 Frakes, W. B. and Succi, G. (2001). An industrial study of reuse, quality, and productivity. *Journal of*
 1017 *Systems and Software* 57, 99–106
- 1018 Fu, R., Qin, X., Dardha, O., and Steuwer, M. (2021). Row-polymorphic types for strategic rewriting. *CoRR*
 1019 abs/2103.13390
- 1020 Gibbons, J. (2006). Datatype-generic programming. In *Proceedings of the 2006 International Conference*
 1021 *on Datatype-Generic Programming* (Berlin, Heidelberg: Springer-Verlag), SSDGP’06, 1–71

- Goguen, J. A. and Burstall, R. M. (1984). Introducing institutions. In *Logics of Programs*, eds. E. Clarke and D. Kozen (Berlin, Heidelberg: Springer Berlin Heidelberg), 221–256
- Gregor, D., Järvi, J., Siek, J., Stroustrup, B., Reis, G. D., and Lumsdaine, A. (2006). Concepts: linguistic support for generic programming in C++. *OOPSLA '06: Proceedings of the 21st annual ACM SIGPLAN conference on Object-oriented programming systems, languages, and applications*, 291–310doi:http://doi.acm.org/10.1145/1167473.1167499
- Hagedorn, B., Lenfers, J., Kunddefinedhler, T., Qin, X., Gorlatch, S., and Steuwer, M. (2020). Achieving high-performance the functional way: A functional pearl on expressing high-performance optimizations as rewrite strategies. *Proc. ACM Program. Lang.* 4. doi:10.1145/3408974
- Hagedorn, B., Stoltzfus, L., Steuwer, M., Gorlatch, S., and Dubach, C. (2018). High performance stencil code generation with lift. In *Proceedings of the 2018 International Symposium on Code Generation and Optimization* (New York, NY, USA: ACM), CGO 2018, 100–112. doi:10.1145/3168824
- Hamre, H.-C. (2022). *Automated Verifications for Magnolia Satisfactions*. Master's thesis, The University of Bergen
- Haveraaen, M. and Roggenbach, M. (2020). Specifying with syntactic theory functors. *Journal of Logical and Algebraic Methods in Programming* 113, 100543. doi:https://doi.org/10.1016/j.jlamp.2020.100543
- Kirchner, H. (2015). *Rewriting Strategies and Strategic Rewrite Programs* (Cham: Springer International Publishing). 380–403. doi:10.1007/978-3-319-23165-5_18
- Lattner, C., Amini, M., Bondhugula, U., Cohen, A., Davis, A., Pienaar, J., et al. (2021). Mlir: Scaling compiler infrastructure for domain specific computation. In *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*. 2–14. doi:10.1109/CGO51591.2021.9370308
- Lemire, D., Kaser, O., and Kurz, N. (2019). Faster remainder by direct computation: Applications to compilers and software libraries. *Softw., Pract. Exper.* 49, 953–970. doi:10.1002/spe.2689
- Liu, A., Bernstein, G. L., Chlipala, A., and Ragan-Kelley, J. (2022). Verified tensor-program optimization via high-level scheduling rewrites. *Proc. ACM Program. Lang.* 6. doi:10.1145/3498717
- Martí-Oliet, N., Meseguer, J., and Verdejo, A. (2005). Towards a strategy language for maude. *Electron. Notes Theor. Comput. Sci.* 117, 417–441
- Martí-Oliet, N., Meseguer, J., and Verdejo, A. (2009). A rewriting semantics for maude strategies. *Electron. Notes Theor. Comput. Sci.* 238, 227–247. doi:10.1016/j.entcs.2009.05.022
- Mullin, L. (1988). *A Mathematics of Arrays*. Ph.D. thesis, Syracuse University
- Mullin, L. and Thibault, S. (1994). *Reduction Semantics for Array Expressions: The Psi Compiler*. Tech. Rep. CSC 94-05, Dept. of CS, University of Missouri-Rolla
- Mullin, L. M. R. and Jenkins, M. A. (1996). Effective data parallel computation using the psi calculus. *Concurrency - Practice and Experience* 8, 499–515. doi:10.1002/(SICI)1096-9128(199609)8:7<499::AID-CPE230>3.0.CO;2-1
- Musser, D. R. and Stepanov, A. A. (1988). Generic programming. In *Symbolic and Algebraic Computation, International Symposium ISSAC'88, Rome, Italy, July 4-8, 1988, Proceedings*, ed. P. M. Gianni (Springer), vol. 358 of *Lecture Notes in Computer Science*, 13–25. doi:10.1007/3-540-51084-2_2
- Nazareth, D. L. and Rothenberger, M. A. (2004). Assessing the cost-effectiveness of software reuse: A model for planned reuse. *Journal of Systems and Software* 73, 245–255
- Puschel, M., Moura, J., Johnson, J., Padua, D., Veloso, M., Singer, B., et al. (2005). Spiral: Code generation for dsp transforms. *Proceedings of the IEEE* 93, 232–275. doi:10.1109/JPROC.2004.840306
- Ragan-Kelley, J., Adams, A., Paris, S., Levoy, M., Amarasinghe, S., and Durand, F. (2012). Decoupling algorithms from schedules for easy optimization of image processing pipelines. *ACM Trans. Graph.* 31. doi:10.1145/2185520.2185528

- 1067 Sannella, D. and Tarlecki, A. (1996). Mind the gap! Abstract versus concrete models of specifications.
 1068 In *Mathematical Foundations of Computer Science 1996, 21st International Symposium, MFCS'96,*
 1069 *Cracow, Poland, September 2-6, 1996, Proceedings*, eds. W. Penczek and A. Szalas (Springer), vol. 1113
 1070 of *Lecture Notes in Computer Science*, 114–134. doi:10.1007/3-540-61550-4_143
- 1071 Steuwer, M., Fensch, C., Lindley, S., and Dubach, C. (2015). Generating performance portable code using
 1072 rewrite rules: From high-level functional expressions to high-performance opencl code. In *Proceedings*
 1073 *of the 20th ACM SIGPLAN International Conference on Functional Programming* (New York, NY, USA:
 1074 Association for Computing Machinery), ICFP 2015, 205–217. doi:10.1145/2784731.2784754
- 1075 Steuwer, M., Remmelg, T., and Dubach, C. (2017). Lift: A functional data-parallel ir for high-performance
 1076 gpu code generation. In *Proceedings of the 2017 International Symposium on Code Generation and*
 1077 *Optimization* (IEEE Press), CGO '17, 74–85. doi:10.1109/CGO.2017.7863730
- 1078 Tang, X. and Järvi, J. (2015). Axioms as generic rewrite rules in C++ with concepts. *Science of*
 1079 *Computer Programming* 97, 320–330. doi:https://doi.org/10.1016/j.scico.2014.05.006. Object-Oriented
 1080 Programming and Systems (OOPS 2010) Modeling and Analysis of Compositional Software (papers
 1081 from EUROMICRO SEAA'12)
- 1082 Vasilache, N., Zinenko, O., Bik, A. J. C., Ravishankar, M., Raoux, T., Belyaev, A., et al. (2022).
 1083 Composable and modular code generation in mlir: A structured and retargetable approach to tensor
 1084 compiler construction doi:10.48550/ARXIV.2202.03293
- 1085 Visser, E. (2005). A survey of strategies in rule-based program transformation systems. *Journal of*
 1086 *Symbolic Computation* 40, 831–873. doi:https://doi.org/10.1016/j.jsc.2004.12.011. Reduction Strategies
 1087 in Rewriting and Programming special issue
- 1088 Wolfe, M. (2021). Performant, portable, and productive parallel programming with standard languages.
 1089 *Computing in Science Engineering* 23, 39–45. doi:10.1109/MCSE.2021.3097167

6.2 A CUDA implementation

To showcase the versatility of the approach to generic code transformations and PDE solvers given in section 6.1, a couple of different versions of the solver were implemented. Among these were OpenMP and CUDA [36] implementations, to highlight how one could parallelize key parts of the algorithm to achieve a significant decrease in computing time on different hardware. In this section we take a closer look at two iterations of the solver implemented in CUDA C++ to leverage the SIMD parallel processing approach utilized by Nvidias line of GPUs.

CUDA is an API developed by Nvidia to facilitate general purpose programming on its lines of GPUs. It is designed to work with high-performance languages such as C, C++ and Fortran. The GPU does not share most of its memory with the CPU, and as such we need to introduce a separation between CPU(host) code and GPU(device) code. In CUDA C++ this is achieved using annotations to mark function as either callable from CPU only, GPU only or both.

```
1 // executes on the CPU
2 __host__ int addTwo_cpu(int a);
3 // executes on the GPU
4 __device__ int addTwo_gpu(int a);
5 // can be executed on both CPU and GPU
6 __host__ __device__ int addTwo_gpu_cpu(int a);
7 // entry point from CPU to GPU
8 __global__ void kernel();
```

Listing 6.1: CUDA Annotations

Program execution starts on the CPU, and in order to access the device side code we need a way to call device side code from the CPU. This is done via global functions, also called kernels. Functions annotated as global can be invoked from the CPU and run device side code. In kernel calls you specify the number of blocks of memory and threads per block the GPU has available for computation. Since CUDA 5.0, device side kernel launches are allowed, paving the way for multiple layers of non-uniform device side memory allocation.

Implementation ¹

The backend for the PDE solver presented in section 6.1 is implemented in C++. In the standard pipeline, magnoliac generates C++ code for the defined Magnolia concepts and

¹The code for both implementations are publicly available online [28, 29].

programs, then compiles the generated code and the user-provided backend using GCC `g++`. CUDA-annotated C++ code is not compatible with ISO standard C++, and as such we are in need of a compiler with CUDA support. For this implementation Nvidias proprietary NVCC was chosen. It is also a feature-complete C++ compiler up to C++17, so for our purposes we replaced `g++` with NVCC for compilation of all C++ code.

Results

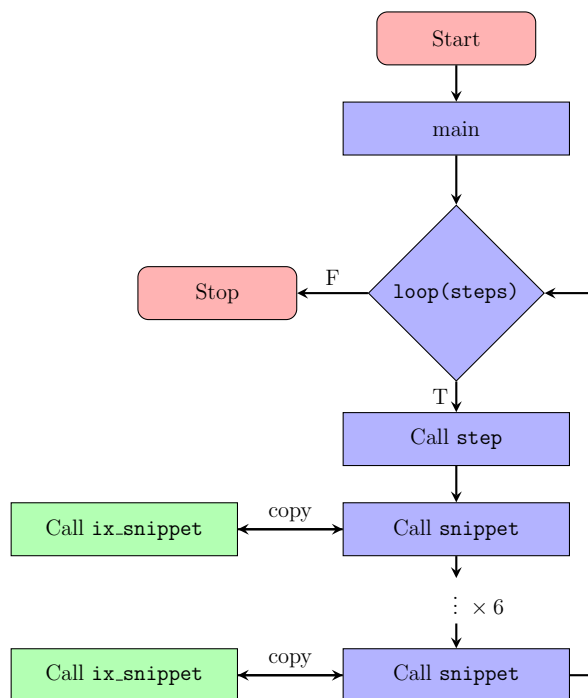


Figure 6.1: Dataflow between CPU(blue) and GPU(green), 1st iteration of the solver.

The first iteration of runtime results fell short of expectations. One would expect much faster speeds from a cutting-edge GPU like the Nvidia A100. The times in Table 6.1 are after 10 steps of the solver, whereas the benchmarks in section 6.1 are after 50 steps.

real	1m12.703s
usr	0m36.333s
sys	0m34.199s

Table 6.1: 10 steps of the CUDA PDE solver with array dimensions 512^3 , ran on a Nvidia Volta A100/80GB. Timed in bash with `time`.

Profiling the binary produced by NVCC, we identified the main causes for the bad performance.

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Name	This is
45.8	21931544537	840	26108981.6	cudaMemcpy	
38.7	18537571949	720	25746627.7	cudaMalloc	
15.4	7377761005	60	32944.5	cudaDeviceReset	

highly ineffective, to stay on the GPU as long as possible

Table 6.2: Snippet of `gpumemtimesum` result generated from `nsys` profile `<pde.bin>`.

We see that calls to `cudaMemcpy` and `cudaMalloc` account for over 80% of the execution time. CUDA memory allocations and copies are expensive operations. In the first iteration of the implementation, memory is allocated and copied between host and device memory 6 times per solver step. Figure 6.1 depicts the dataflow between CPU and GPU in the first iteration of the CUDA implementation.

Improvements

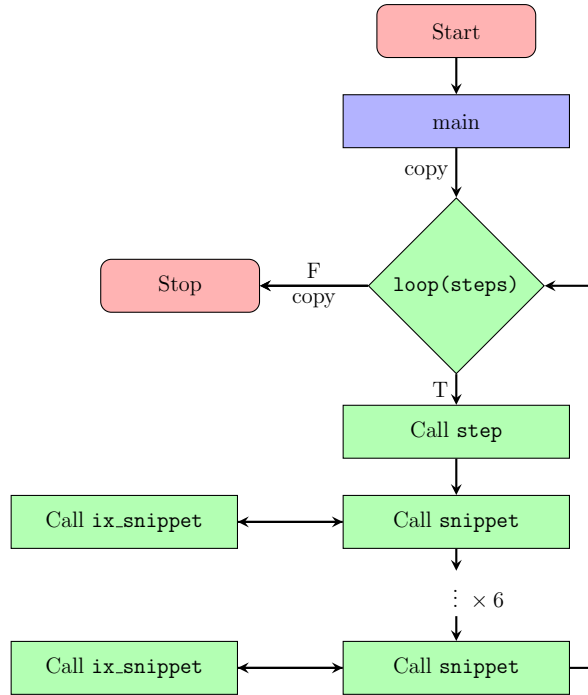


Figure 6.2: Improved dataflow between CPU(blue) and GPU(green)

Following the implementation from section 6.1, we theorized that we could reduce the number of copies between host and device to a single copy to GPU before solver execution and a single copy to CPU after, resulting in a significant reduction in execution time. Due to time constraints, this improved implementation was not completed before the paper

submission deadline. Figure 6.2 depicts the theorized improved dataflow between GPU and CPU.

Remark on earlier work

The work presented in Burrows et al. includes runtime experiments for two CUDA versions implemented in the process. Analyzing the source code for these implementation provided us with valuable insights in how to approach the the problem. An important difference between the approach presented in Burrows et al. and the the one presented in this thesis is the density of the kernel computations. Previous implementations opted for implementing each operation as separate kernel calls, resulting in low density of computation. In other words, quite large parts of available GPU cores remain unused per computation. With this in mind, the approach taken in section 6.1 relies on a dense, inlined **step** function presented in Chetioui et al. With proper memory management, the theorized improvement depicted in Figure 6.2 should leverage the cores available to us on the GPU to a larger degree.

Chapter 7

Future Work & Conclusion

7.1 Future Work

Even though the thesis process draws to an end, there are multiple places to start if one would wish to continue the work discussed here. An natural first step would be to continue the work on the MoA API. Improving the existing code should be prioritized, as the current state of the code base suffers from a focus on producing compilable code on a tight deadline. In particular, this meant sacrificing time that could have been spent fleshing out the array specification. The API in its current form is also prioritizing a limited subset of the theory to explore the specific API proposed by Burrows et al. Continuing to investigate the parts of the ψ -calculus not covered in this thesis would come as a natural continuation. I believe that a logical first extension of the API would be to fully specify and implement padding, taking it a step closer to its intended state of usability. One could also move towards extending it to coverONF, both specified in the generic API and as a step toward observing how implementations on different hardware may differ.

Although developed independently of existing Magnolia library packages, once in an acceptable state, integration of the API with the standard library may be useful for future Magnolia projects.

Completion of the improved CUDA implementation discussed in section 6.2 could add weight to the argument presented in section 6.1 that generating performant code from generic languages like Magnolia can compete with existing compiler tools. Chetioui et al. have already demonstrated promising results for this approach in another domain.

7.2 Conclusion

In this thesis we have explored the MoA calculus and how it can serve as a foundation for generic multiarrays. We have presented relevant background theory, creating an arena for discussion around generic programming and API design. We have showed that one can abstract hardware specific details away, without losing the ability to target specific architectures. Implementing a subset of the MoA theory in Magnolia, we have gained knowledge about how we can leverage the Magnolia type system to go from generic specifications and concrete implementations while retaining type safety.

As hardware improves, it is critical to keep software portable and maintainable while retaining performance. This principle is followed by leveraging concepts from generic programming to abstract away from hardware specific implementations. We have seen that Magnolia allows us to reason about programs at a higher level, capturing this philosophy. The modularity gained by designing software with a clear separation between specification and implementation also serves the purpose of forcing developers away from the typical monolithic structure of large systems, where dependencies are often difficult to swap out. Haverdaen argues that allocating more time to domain exploration is a cost effective way to design software long term, and we have been given a taste of this through the lens of re-use mechanisms in Magnolia such as renamings. We hope to see a gradual shift in the software design philosophy, towards a more modular future for programming.

Glossary

ψ -reduction The process of transforming an array expression into an equivalent expression only using the ψ operator.

AEP Annual Energy Production.

API Application Programming Interface.

APL A Programming Language (APL), an array programming language developed in the 1960s.

BLDL Bergen Language Design Laboratory.

CPU Central Processing Unit.

CUDA Parallel computing platform developed by Nvidia for their line of GPUs.

DNF Denotational Normal Form.

FLOPS Floating Point Operations Per Second.

GCC GNU Compiler Collection.

GPGPU General-purpose Programming on Graphical Processing Units.

GPU Graphical Processing Unit.

HPC High-performance computing.

Magnolia Magnolia is a programming language based on the theory of institutions.

magnoliac A Magnolia compiler under active development at BLDL.

MoA A Mathematics of Arrays.

MPI Message Passing Interface.

NVCC Nvidia CUDA Compiler.

ONF Operational Normal Form.

OpenBLAS Open Basic Linear Algebra Subprograms.

OpenCL Open Computing Language.

OpenMP Open Multi-Processing, shared-memory multiprocessing programming API.

PDE Partial Differential Equation.

SIMD Single Instruction, Multiple Data.

SMT Satisfiability Modulo Theories.

VLSI Very Large Scale Integration.

Bibliography

- [1] Ole Jørgen Abusdal. Transformations for array programming. Master’s thesis, University of Bergen, 2020.
- [2] Anya Helene Bagge. *Constructs & Concepts: Language Design for Flexibility and Reliability*. PhD thesis, Research School in Information and Communication Technology, Department of Informatics, University of Bergen, Norway, PB 7803, 5020 Bergen, Norway, 2009.
URL: <http://www.ii.uib.no/~anya/phd/>.
- [3] Eva Burrows, Helmer André Friis, and Magne Haveraaen. An array api for finite difference methods. In *roceedings of the 5th ACM SIGPLAN International Workshop on Libraries, Languages, and Compilers for Array Programming*. Association for Computing Machinery, 2018. doi: 10.1145/3219753.3219761.
URL: <https://doi.org/10.1145/3219753.3219761>. 59–66.
- [4] J. Carlson, J.A.C.A.J.A. Wiles, J.A. Carlson, A. Jaffe, A. Wiles, Clay Mathematics Institute, and American Mathematical Society. *The Millennium Prize Problems*. American Mathematical Society, 2006. ISBN 9780821836798.
- [5] Benjamin Chetoui. magnoliac: A magnolia compiler, December 2020.
- [6] Benjamin Chetoui, Lenore Mullin, Ole Abusdal, Magne Haveraaen, Jaakko Järvi, and Sandra Macià. Finite difference methods fengshui: alignment through a mathematics of arrays. In *Proceedings of the 6th ACM SIGPLAN International Workshop on Libraries, Languages and Compilers for Array Programming*. Association for Computing Machinery, 2019. doi: 10.1145/3315454.3329954.
URL: <https://doi.org/10.1145/3315454.3329954>. 2–13.
- [7] Benjamin Chetoui, Ole Abusdal, Magne Haveraaen, Jaakko Järvi, and Lenore Mullin. Padding in the mathematics of arrays. In *Proceedings of the 7th ACM SIGPLAN International Workshop on Libraries, Languages and Compilers for Array*

- Programming*. Association for Computing Machinery, 2021. doi: 10.1145/3460944.3464311.
URL: <https://doi.org/10.1145/3460944.3464311>. 15–26.
- [8] Benjamin Chetoui, Jaakko Järvi, and Magne Haveræen. Revisiting language support for generic programming: The case of magnolia. Under review, 2022.
 - [9] Benjamin Chetoui, Marius Kleppe Larnøy, Jaakko Järvi, Magne Haveræen, and Lenore Mullin. P³ problem and magnolia language: Specializing array computations for emerging architectures. Under review, 2022.
 - [10] James Dehnert and Alexander Stepanov. Fundamentals of generic programming. volume 1766, pages 1–11, 01 1998. ISBN 978-3-540-41090-4. doi: 10.1007/3-540-39953-4_1.
 - [11] Python Software Foundation. array - efficient arrays of numeric values.
URL: <https://docs.python.org/3/library/array.html>. [Online; accessed 31.05.22].
 - [12] Jeremy Gibbons. Datatype-generic programming. In *Proceedings of the 2006 International Conference on Datatype-Generic Programming*, SSDGP’06, page 1–71, Berlin, Heidelberg, 2006. Springer-Verlag. ISBN 3540767851.
 - [13] Joseph A. Goguen and Rod M. Burstall. Institutions: Abstract model theory for specification and programming. *J. ACM*, 39(1):95–146, January 1992. ISSN 0004-5411. doi: 10.1145/147508.147524.
URL: <https://doi.org/10.1145/147508.147524>.
 - [14] John L. Gustafson. *Moore’s Law*, pages 1177–1184. Springer US, Boston, MA, 2011. ISBN 978-0-387-09766-4. doi: 10.1007/978-0-387-09766-4_81.
URL: https://doi.org/10.1007/978-0-387-09766-4_81.
 - [15] John L. Gustafson and Lenore M. Mullin. Tensors come of age: Why the ai revolution will help hpc, 2017.
 - [16] G. Hains and L. M. R. Mullin. Parallel functional programming with arrays. *The Computer Journal*, 36(3):238–245, 01 1993. ISSN 0010-4620. doi: 10.1093/comjnl/36.3.238.
URL: <https://doi.org/10.1093/comjnl/36.3.238>.
 - [17] Hans-Christian Hamre. Automated verifications for magnolia satisfactions. Master’s thesis, The University of Bergen, 2022.

- [18] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerkwijk, Matthew Brett, Allan Haldane, Jaime Fernández del Río, Mark Wiebe, Pearu Peterson, Pierre Gérard-Marchant, Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array programming with NumPy. *Nature*, 585(7825):357–362, September 2020. doi: 10.1038/s41586-020-2649-2.
URL: <https://doi.org/10.1038/s41586-020-2649-2>.
- [19] Kristoffer Haugsbakk. Program transformations in magnolia. Master’s thesis, The University of Bergen, 2017.
- [20] Magne Haverlaen. Domain engineering the magnolia way. In Alexander K. Petrenko and Andrei Voronkov, editors, *Perspectives of System Informatics*, pages 196–210, Cham, 2018. Springer International Publishing. ISBN 978-3-319-74313-4.
- [21] S. Hoyer and J. Hamman. xarray: N-D labeled arrays and datasets in Python. *Journal of Open Research Software*, 5(1), 2017. doi: 10.5334/jors.148.
URL: <https://doi.org/10.5334/jors.148>.
- [22] J. D. Hunter. Matplotlib: A 2d graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007. doi: 10.1109/MCSE.2007.55.
- [23] ISO. *Draft International Standard ISO/IEC 1539-1:2004(E): Information technology — Programming languages — Fortran Part 1: Base Language*. May 2004.
URL: <https://wg5-fortran.org/N1601-N1650/N1601.pdf>.
- [24] ISO. *ISO/IEC 9899:2011 Information technology — Programming languages — C*. December 2011.
URL: http://www.iso.org/iso/iso_catalogue/catalogue_tc/catalogue_detail.htm?csnumber=57853.
- [25] Kenneth E. Iverson. *A Programming Language*. John Wiley & Sons, Inc., USA, 1962. ISBN 0471430145.
- [26] Neo Shih-Chao Kao and Tony Wen-Hann Sheu. Development of a finite element flow solver for solving three-dimensional incompressible navier–stokes solutions on multiple gpu cards. *Computers & Fluids*, 167:285–291, 2018. ISSN 0045-7930. doi: <https://doi.org/10.1016/j.compfluid.2018.03.033>.

- [27] I G W Krabben, M P van der Laan, M. Koivisto, T J Larsen, M M Pedersen, and K S Hansen. Why curved wind turbine rows are better than straight ones. *Journal of Physics: Conference Series*, 1256(1):012028, jul 2019. doi: 10.1088/1742-6596/1256/1/012028.
URL: <https://doi.org/10.1088/1742-6596/1256/1/012028>.
- [28] Marius Larnøy. mariuslarnoy/magnolia-lang at cuda-dynamic, .
URL: <https://github.com/mariuslarnoy/magnolia-lang/tree/cuda-dynamic>. [Online; accessed 03.05.22].
- [29] Marius Larnøy. mariuslarnoy/magnolia-lang at cuda, .
URL: <https://github.com/mariuslarnoy/magnolia-lang/tree/cuda>. [Online; accessed 03.05.22].
- [30] Marius Larnøy. mariuslarnoy/magnolia-lang at naturalnumbers, .
URL: <https://github.com/mariuslarnoy/magnolia-lang/tree/naturalnumbers>. [Online; accessed 25.05.22].
- [31] Marius Larnøy. mariuslarnoy/magnolia-lang at moa-ndim, .
URL: <https://github.com/mariuslarnoy/magnolia-lang/tree/moa-ndim>. [Online; accessed 20.05.22].
- [32] Lars Moastuen. Real-time simulation of the incompressible navier-stokes equations on the gpu. Master’s thesis, The University of Oslo, 2007.
- [33] Lenore Mullin and Scott Thibault. A reduction semantics for array expressions: The psi compiler. 1994.
- [34] Preferred Networks. Cupy.
URL: <https://cupy.dev/>. [Online; accessed 12.05.22].
- [35] NumPy. Case study: First image of a black hole.
URL: <https://numpy.org/case-studies/blackhole-image/>. [Online; accessed 12.05.22].
- [36] NVIDIA, Péter Vingelmann, and Frank H.P. Fitzek. Cuda, release: 10.2.89, 2020.
URL: <https://developer.nvidia.com/cuda-toolkit>. [Online; accessed 03.05.22].
- [37] C. Ostrouchov and L. Mullin. python-moa, 2019.
URL: <https://github.com/Quansight-Labs/python-moa>. [Online; accessed 03.05.22].
- [38] Mads M. Pedersen, Paul van der Laan, Mikkel Friis-Møller, Jennifer Rinker, and Pierre-Elouan Réthoré. Dtuwindenergy/pywake: Pywake. Feb 2019. doi: 10.5281/zenodo.2562662.

- [39] Wileam Phan, Jeremie Vandenplas, Arjen Markus, and Lenore Mullin. Mathematics of arrays library for modern fortran, 2021.
URL: <https://github.com/wyphan/moa-fortran>. [Online; accessed 03.05.22].
- [40] Lenore Marie Restifo Mullin. *A Mathematics of Arrays*. Thesis, Syracuse University, 1988.
- [41] Riccardo Riva, Jaime Liew, Mikkel Friis-Møller, Nikolay Dimitrov, Emre Barlas, Pierre-Elouan Réthoré, and Arvydas Beržonskis. Wind farm layout optimization with load constraints using surrogate modelling. *Journal of Physics: Conference Series*, 1618(4):042035, sep 2020. doi: 10.1088/1742-6596/1618/4/042035.
URL: <https://doi.org/10.1088/1742-6596/1618/4/042035>.
- [42] Jeremy Siek, Lie-Quan Lee, and Andrew Lumsdaine. *The Boost Graph Library: User Guide and Reference Manual*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2002. ISBN 0-201-72914-8.
- [43] John E. Stone, David Gohara, and Guochun Shi. Opencl: A parallel programming standard for heterogeneous computing systems. *Computing in Science & Engineering*, 12(3):66–73, 2010. doi: 10.1109/MCSE.2010.69.
- [44] M P van der Laan, S J Andersen, and P-E Réthoré. Brief communication: Wind speed independent actuator disk control for faster aep calculations of wind farms using cfd. 2019. doi: 10.5194/wes-2019-27.
- [45] Paul Veers. Three-dimensional wind simulation. Sandia National Laboratories, 01 1988.
- [46] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, Matt Haberland, Tyler Reddy, David Cournapeau, Evgeni Burovski, Pearu Peterson, Warren Weckesser, Jonathan Bright, Stéfan J. van der Walt, Matthew Brett, Joshua Wilson, K. Jarrod Millman, Nikolay Mayorov, Andrew R. J. Nelson, Eric Jones, Robert Kern, Eric Larson, C. J. Carey, İlhan Polat, Yu Feng, Eric W. Moore, Jake VanderPlas, Denis Laxalde, Josef Perktold, Robert Cimrman, Ian Henriksen, E. A. Quintero, Charles R. Harris, Anne M. Archibald, Antônio H. Ribeiro, Fabian Pedregosa, Paul van Mulbregt, Aditya Vijaykumar, Alessandro Pietro Bardelli, Alex Rothberg, Andreas Hilboll, Andreas Kloeckner, Anthony Scopatz, Antony Lee, Ariel Rokem, C. Nathan Woods, Chad Fulton, Charles Masson, Christian Häggström, Clark Fitzgerald, David A. Nicholson, David R. Hagen, Dmitrii V. Pasechnik, Emanuele Olivetti, Eric Martin, Eric Wieser, Fabrice Silva, Felix Lenders, Florian Wilhelm, G. Young, Gavin A. Price, Gert-Ludwig Ingold, Gregory E. Allen, Gregory R. Lee, Hervé Audren, Irvin

Probst, Jörg P. Dietrich, Jacob Silterra, James T. Webber, Janko Slavič, Joel Nothman, Johannes Buchner, Johannes Kulick, Johannes L. Schönberger, José Vinícius de Miranda Cardoso, Joscha Reimer, Joseph Harrington, Juan Luis Cano Rodríguez, Juan Nunez-Iglesias, Justin Kuczynski, Kevin Tritz, Martin Thoma, Matthew Newville, Matthias Kümmerer, Maximilian Bolingbroke, Michael Tartre, Mikhail Pak, Nathaniel J. Smith, Nikolai Nowaczyk, Nikolay Shebanov, Oleksandr Pavlyk, Per A. Brodtkorb, Perry Lee, Robert T. McGibbon, Roman Feldbauer, Sam Lewis, Sam Tygier, Scott Sievert, Sebastiano Vigna, Stefan Peterson, Surhud More, Tadeusz Pudlik, Takuya Oshima, Thomas J. Pingel, Thomas P. Robitaille, Thomas Spura, Thouis R. Jones, Tim Cera, Tim Leslie, Tiziano Zito, Tom Krauss, Utkarsh Upadhyay, Yaroslav O. Halchenko, Yoshiki Vázquez-Baeza, and SciPy 1.0 Contributors. Scipy 1.0: fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3):261–272, Mar 2020. ISSN 1548-7105. doi: 10.1038/s41592-019-0686-2.

URL: <https://doi.org/10.1038/s41592-019-0686-2>.

- [47] J.A. Vitulli, G.C. Larsen, M.M. Pedersen, S. Ott, and M. Friis-Møller. Optimal open loop wind farm control. *Journal of Physics: Conference Series*, 1256(1):012027, jul 2019. doi: 10.1088/1742-6596/1256/1/012027.

URL: <https://doi.org/10.1088/1742-6596/1256/1/012027>.

- [48] Wes McKinney. Data Structures for Statistical Computing in Python. In Stéfan van der Walt and Jarrod Millman, editors, *Proceedings of the 9th Python in Science Conference*, pages 56 – 61, 2010. doi: 10.25080/Majora-92bf1922-00a.